

GPU programming with OpenCV Samples

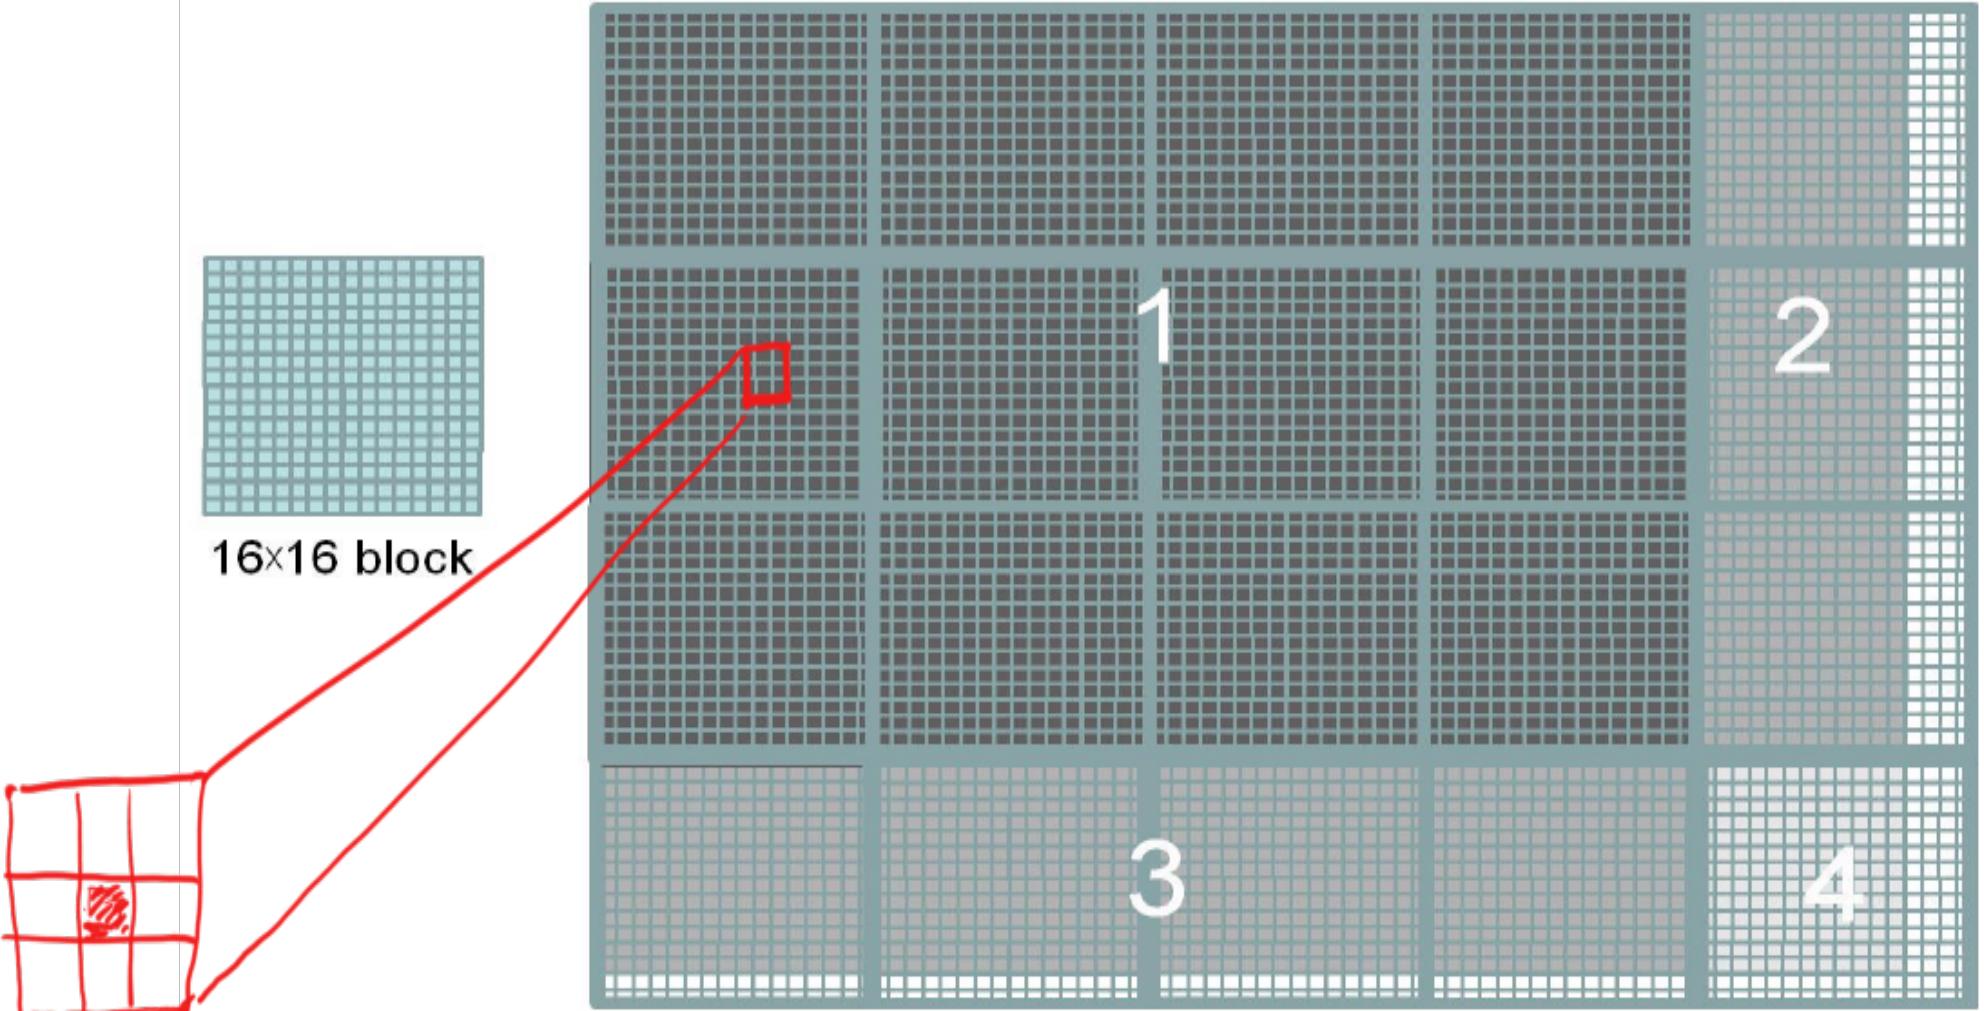
By Mohit Kumar

CUDA:GPU:MultiDim:ex-3(Blur)



Image blurring

CUDA:GPU:MultiDim:ex-3(Blur)



Not all threads in a Block will follow the same control flow path.

CUDA:GPU:MultiDim:ex-3(Blur)

9	1	3	4	8
10	4	6	0	10
2	4	1	2	1
2	3	0	3	1
8	0	0	3	9

Avg Intensity Value =

$$\frac{1}{9} (4 + 10 + 4 + 6 + 2 + 1 + 6 + 2 + 3)$$

$$(w_1 \cdot 4) + (w_2 \cdot 10) + (w_3 \cdot 4) + (w_4 \cdot 6).$$

- weighted sum generates a far smoother image than just taking an average.

Weighted



unweighted.

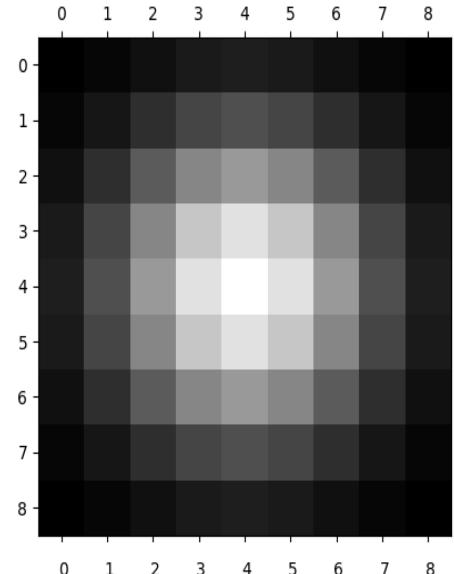


CUDA:GPU:MultiDim:ex-3(Blur)

Visualizing the filter

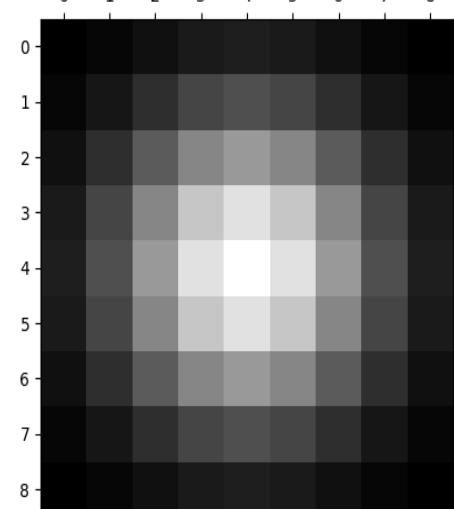
```
state=np.matrix(  
    [ [ 0.018316, 0.043937, 0.082085, 0.119433, 0.135335, 0.119433, 0.082085, 0.043937, 0.018316],  
      [ 0.043937, 0.105399, 0.196912, 0.286505, 0.324652, 0.286505, 0.196912, 0.105399, 0.043937],  
      [ 0.082085, 0.196912, 0.367879, 0.535261, 0.606531, 0.535261, 0.367879, 0.196912, 0.082085],  
      [ 0.119433, 0.286505, 0.535261, 0.778801, 0.882497, 0.778801, 0.535261, 0.286505, 0.119433],  
      [ 0.135335, 0.324652, 0.606531, 0.882497, 1.000000, 0.882497, 0.606531, 0.324652, 0.135335],  
      [ 0.119433, 0.286505, 0.535261, 0.778801, 0.882497, 0.778801, 0.535261, 0.286505, 0.119433],  
      [ 0.082085, 0.196912, 0.367879, 0.535261, 0.606531, 0.535261, 0.367879, 0.196912, 0.082085],  
      [ 0.043937, 0.105399, 0.196912, 0.286505, 0.324652, 0.286505, 0.196912, 0.105399, 0.043937],  
      [ 0.018316, 0.043937, 0.082085, 0.119433, 0.135335, 0.119433, 0.082085, 0.043937, 0.018316]  
    ])  
plt.matshow(state, cmap=plt.cm.gray)  
plt.show()
```

(Non Normalized)



```
state=np.matrix(  
    [ [ 0.000763, 0.001831, 0.003422, 0.004978, 0.005641, 0.004978, 0.003422, 0.001831, 0.000763],  
      [ 0.001831, 0.004393, 0.008208, 0.011942, 0.013532, 0.011942, 0.008208, 0.004393, 0.001831],  
      [ 0.003422, 0.008208, 0.015334, 0.022311, 0.025282, 0.022311, 0.015334, 0.008208, 0.003422],  
      [ 0.004978, 0.011942, 0.022311, 0.032463, 0.036785, 0.032463, 0.022311, 0.011942, 0.004978],  
      [ 0.005641, 0.013532, 0.025282, 0.036785, 0.041683, 0.036785, 0.025282, 0.013532, 0.005641],  
      [ 0.004978, 0.011942, 0.022311, 0.032463, 0.036785, 0.032463, 0.022311, 0.011942, 0.004978],  
      [ 0.003422, 0.008208, 0.015334, 0.022311, 0.025282, 0.022311, 0.015334, 0.008208, 0.003422],  
      [ 0.001831, 0.004393, 0.008208, 0.011942, 0.013532, 0.011942, 0.008208, 0.004393, 0.001831],  
      [ 0.000763, 0.001831, 0.003422, 0.004978, 0.005641, 0.004978, 0.003422, 0.001831, 0.000763]  
    ])  
plt.matshow(state, cmap=plt.cm.gray)  
plt.show()
```

(Normalized)

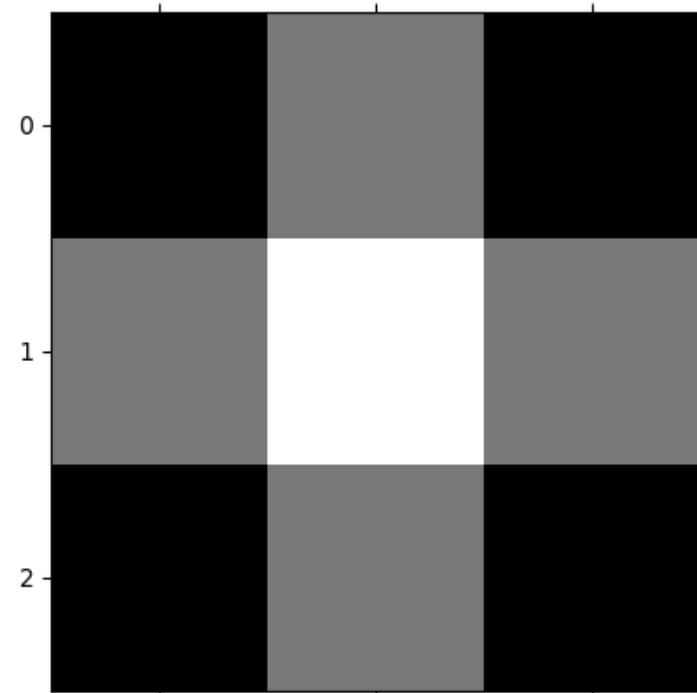
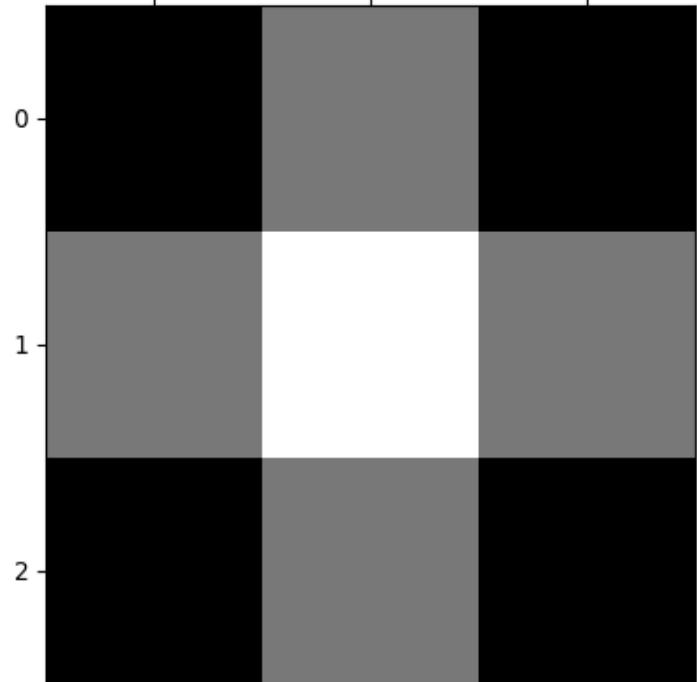


- The tactic with the filter results in a smoother blur than with the average
- Intuitively this happens because the pixel in question has a greater "say" on the resulting pixel than the surrounding pixels.

CUDA:GPU:MultiDim:ex-3(Blur)

```
state=np.matrix(  
    [ [0.778801, 0.882497, 0.778801],  
      [0.882497, 1.000000, 0.882497],  
      [0.778801, 0.882497, 0.778801]  
    ])  
  
plt.matshow(state, cmap=plt.cm.gray)  
plt.show()  
  
state=np.matrix(  
    [ [0.101868, 0.115432, 0.101868],  
      [0.115432, 0.130801, 0.115432],  
      [0.101868, 0.115432, 0.101868]  
    ])  
  
plt.matshow(state, cmap=plt.cm.gray)  
plt.show()
```

- The filter for the slide is (3×3) for easier drawings



CUDA:GPU:MultiDim:ex-3(Blur)

Blur (Logical steps)

1. Separate the RGB channels.
2. Apply the blurring filter to each channel separately
3. Recombine the channel and reconstruct the image.

CUDA:GPU:MultiDim:ex-3(Blur)

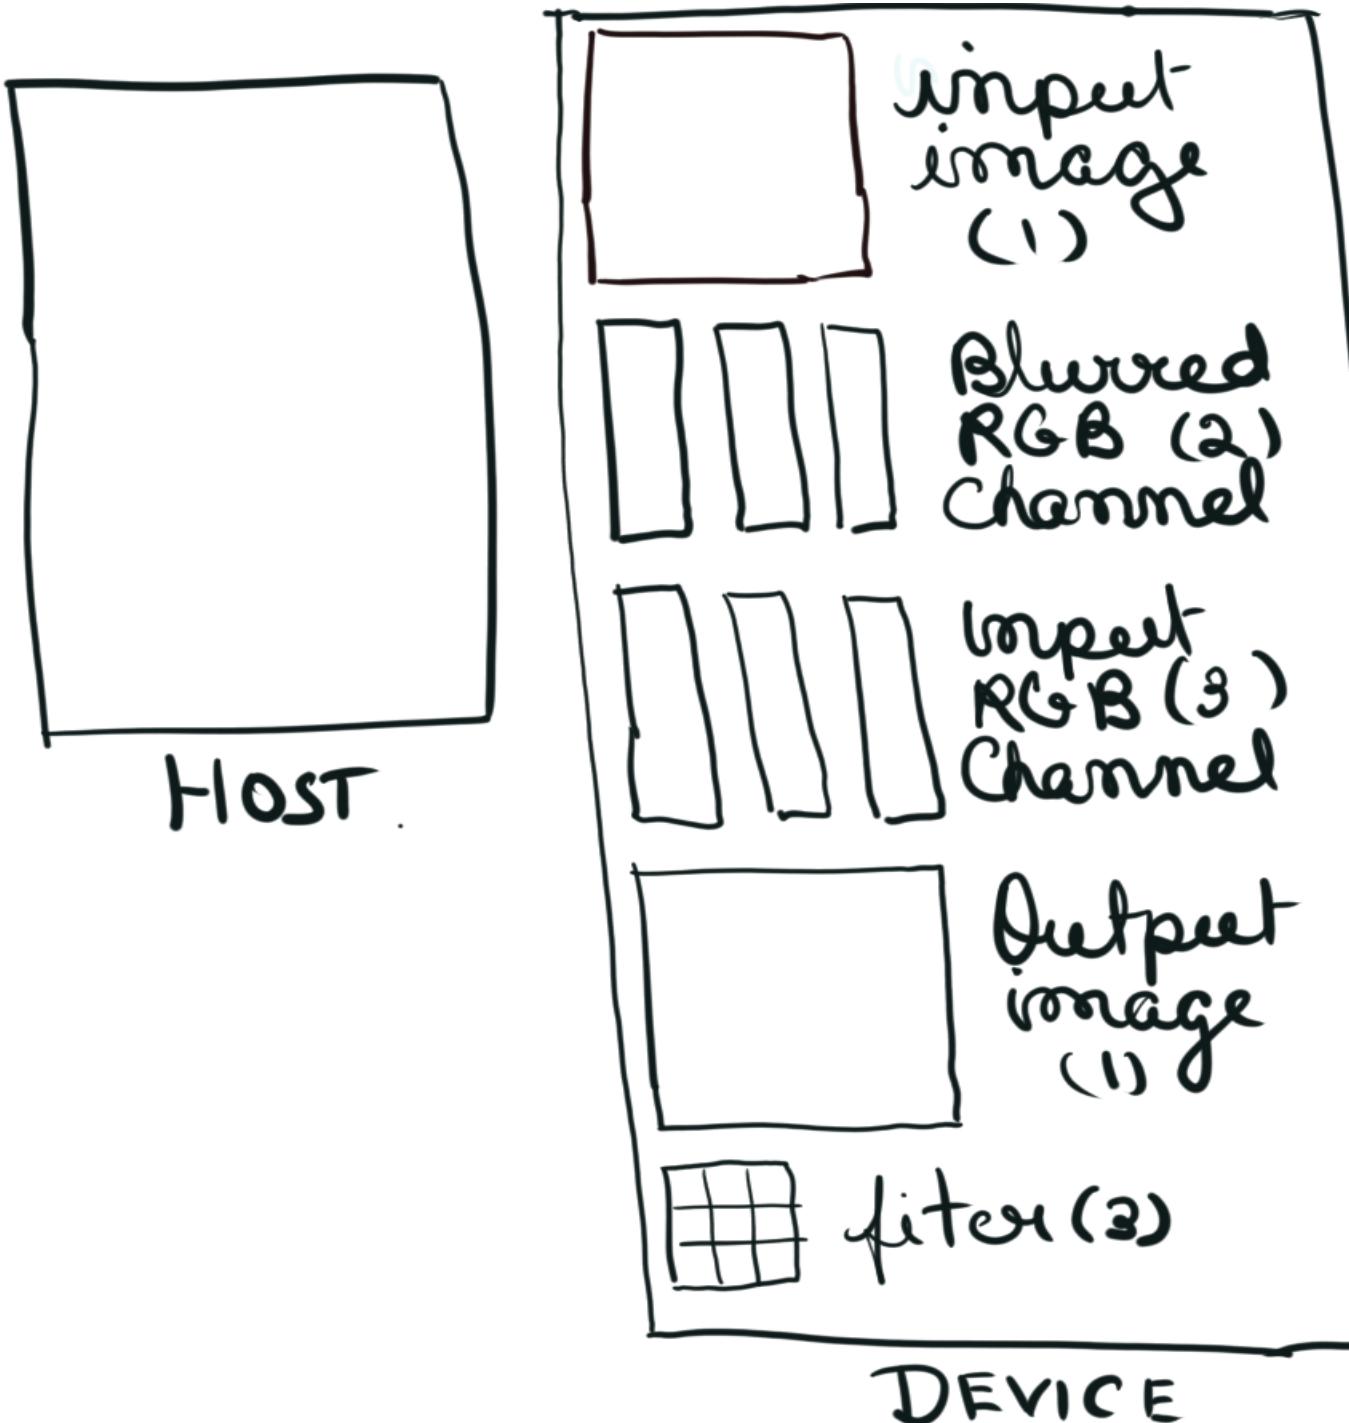
Blur (Actual Steps)

1. Allocate memory for the input and the output images on the device
 - (a) copy image to device
 - (b) memset '0' to output image on device
2. Allocate memory for each channel separately (output) on the device
 - (a) memset all to '0'
3. Allocate memory for each channel separately (input) on the device and for the filter
 - (a) memset RGB memory to '0'
 - (b) copy filter to the device

CUDA:GPU:MultiDim:ex-3(Blur)

4. launch the kernel to separate the RGB
5. separateChannel kernel executes
6. launch the kernel to blur the channels separately
- 6a, 6b, c gaussianBlur executes for all the channels separately
7. Launch the recombine channel
- 7a. recombine kernel executes

CUDA:GPU:MultiDim:ex-3(Blur)



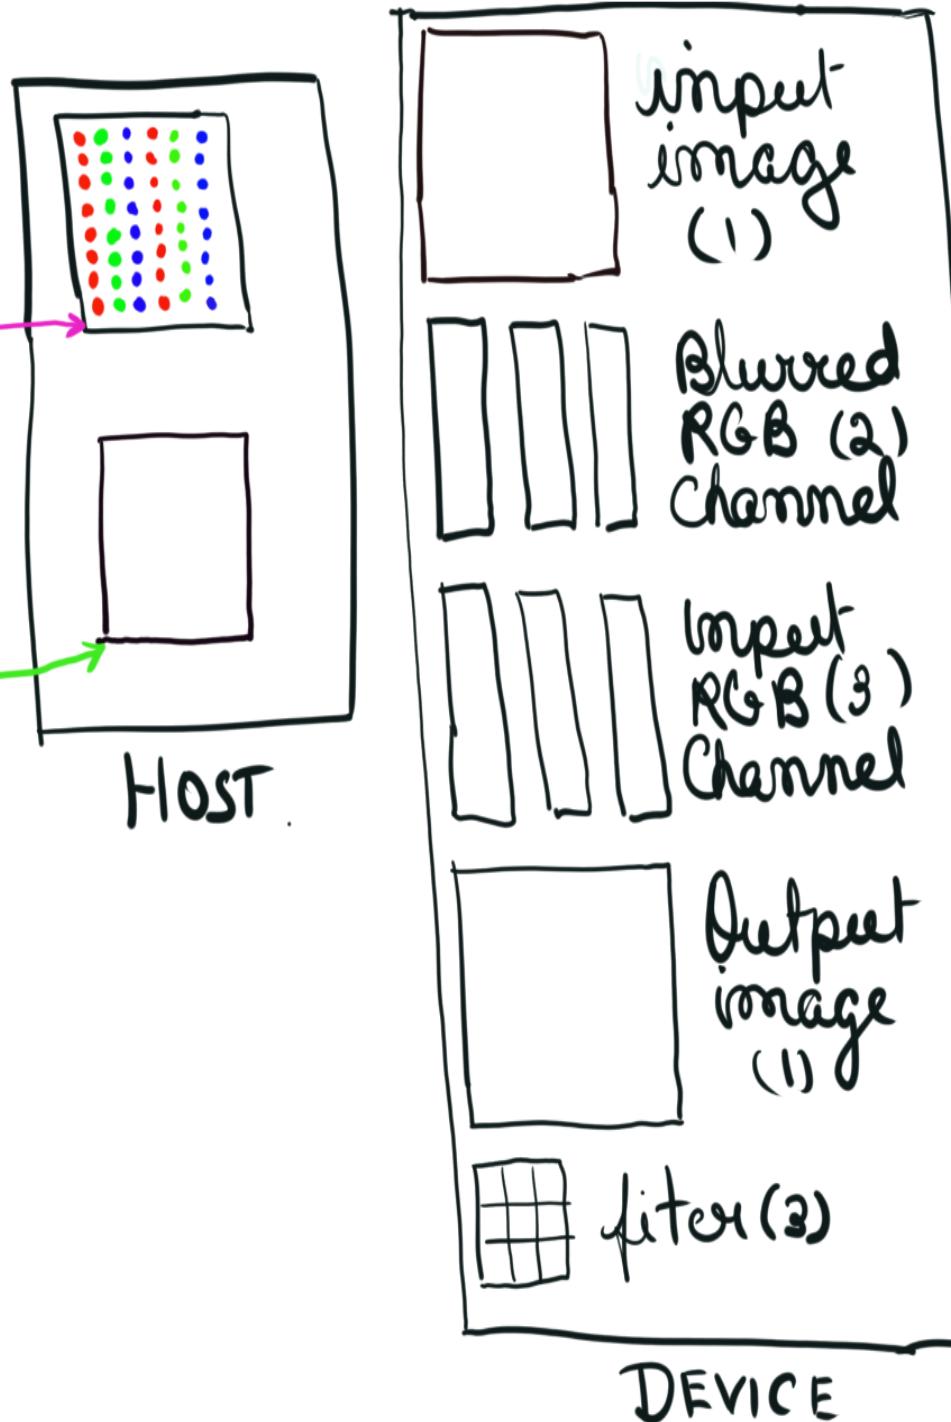
CUDA:GPU:MultiDim:ex-3(Blur)

```
cv::Mat image = cv::imread(filename.c_str(), CV_LOAD_IMAGE_COLOR);
if (image.empty()) {
    std::cerr << "Couldn't open file: " << filename << std::endl;
    exit(1);
}

cv::cvtColor(image, imageInputRGB, CV_BGR2RGB);
//allocate memory for the output
imageOutputRGB.create(image.rows, image.cols, CV_8UC4);

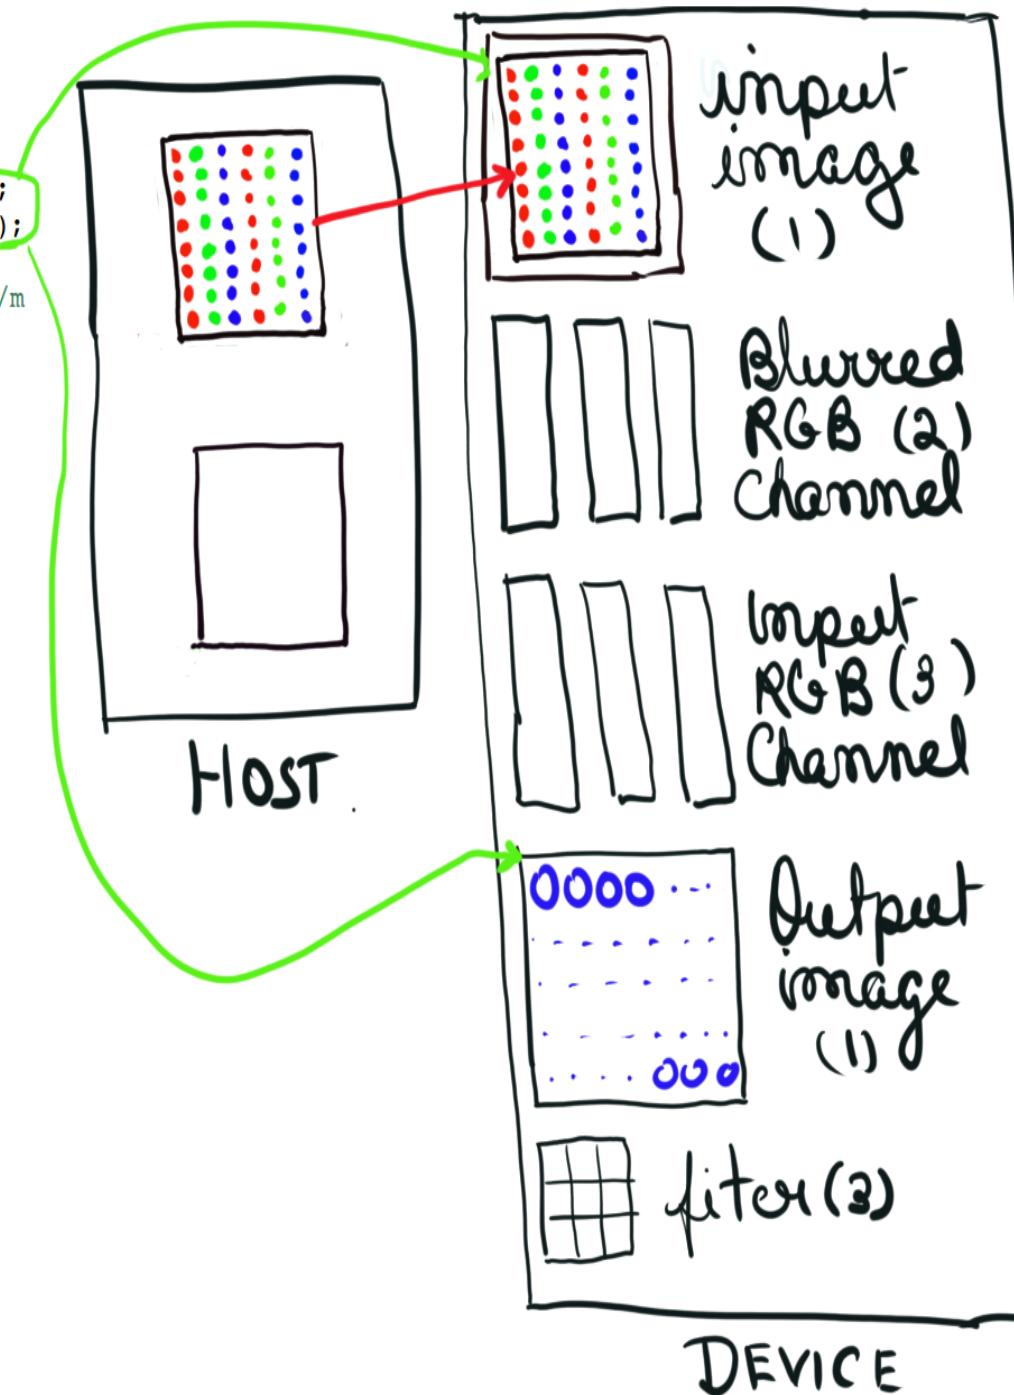
*h_inputImageRGB = (uchar4 *) imageInputRGB.ptr<unsigned char>(0);
*h_outputImageRGB = (uchar4 *) imageOutputRGB.ptr<unsigned char>(0);
```

0. Open the image on the host memory



CUDA:GPU:MultiDim:ex-3(Blur)

```
const size_t numPixels = numRows() * numCols();
//allocate memory on the device for both input and output
checkCudaErrors(cudaMalloc(d_inputImageRGBA, sizeof(uchar4) * numPixels));
checkCudaErrors(cudaMalloc(d_outputImageRGBA, sizeof(uchar4) * numPixels));
checkCudaErrors(
    cudaMemset(*d_outputImageRGBA, 0, numPixels * sizeof(uchar4))), //m
    16
//copy input array to the GPU
checkCudaErrors(
    cudaMemcpy(*d_inputImageRGBA, *h_inputImageRGBA,
        sizeof(uchar4) * numPixels, cudaMemcpyHostToDevice));
    1a
```



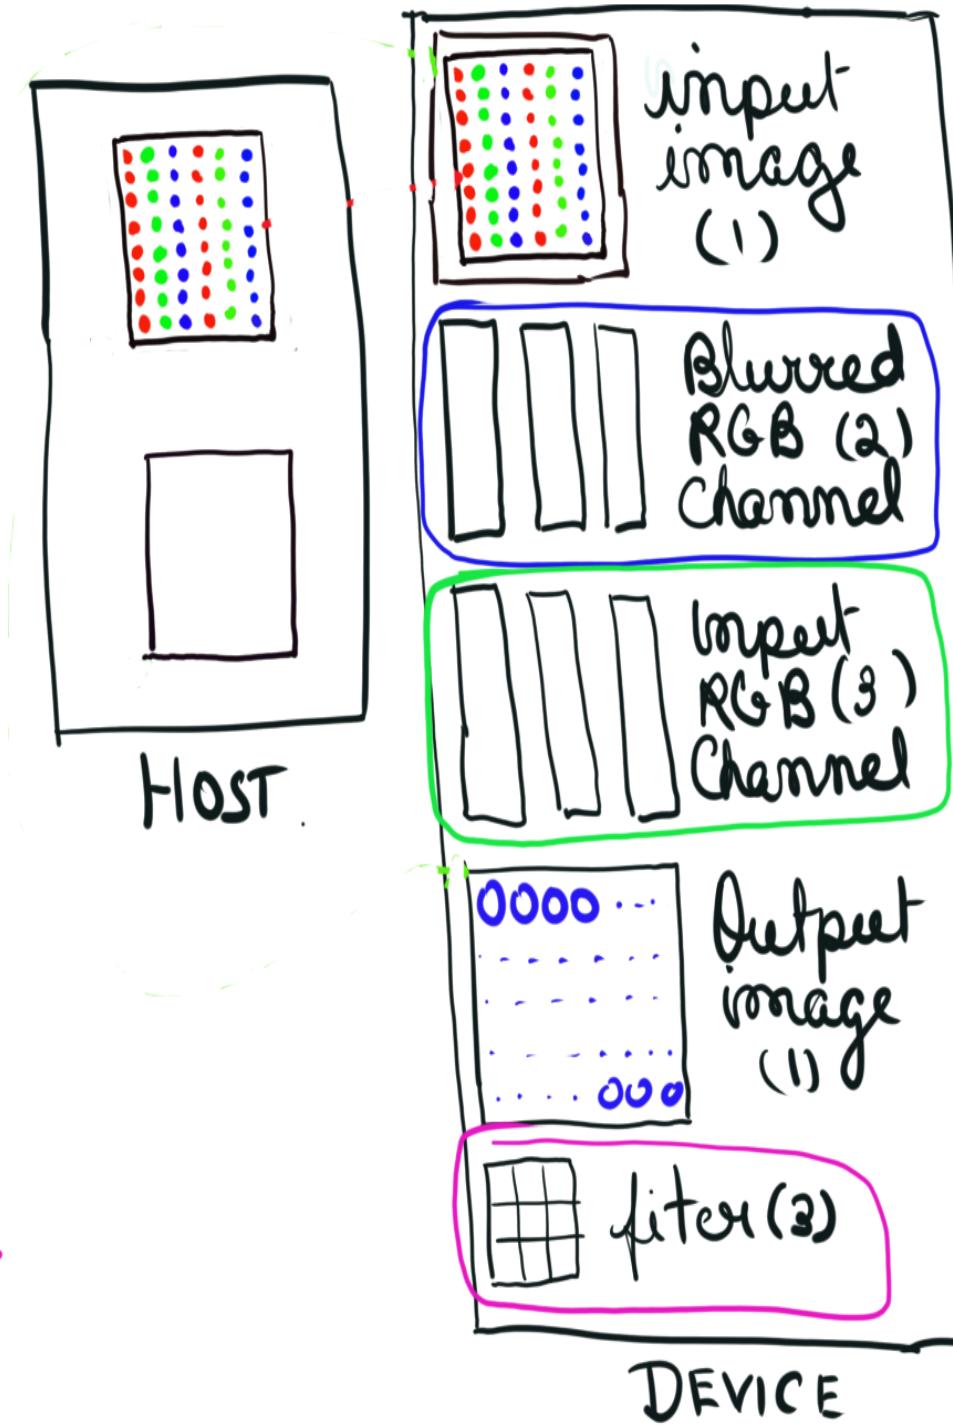
CUDA:GPU:MultiDim:ex-3(Blur)

```
//blurred
checkCudaErrors(
    cudaMalloc(d_redBlurred, sizeof(unsigned char) * numPixels));
checkCudaErrors(
    cudaMalloc(d_greenBlurred, sizeof(unsigned char) * numPixels));
checkCudaErrors(
    cudaMalloc(d_blueBlurred, sizeof(unsigned char) * numPixels));
checkCudaErrors(
    cudaMemcpy(*d_redBlurred, 0, sizeof(unsigned char) * numPixels));
checkCudaErrors(
    cudaMemcpy(*d_greenBlurred, 0, sizeof(unsigned char) * numPixels));
checkCudaErrors(
    cudaMemcpy(*d_blueBlurred, 0, sizeof(unsigned char) * numPixels));

//allocate memory for the three different channels
//original
checkCudaErrors(
    cudaMalloc(&d_red,
              sizeof(unsigned char) * numRowsImage * numColsImage));
checkCudaErrors(
    cudaMalloc(&d_green,
              sizeof(unsigned char) * numRowsImage * numColsImage));
checkCudaErrors(
    cudaMalloc(&d_blue,
              sizeof(unsigned char) * numRowsImage * numColsImage));

size_t filtersize = sizeof(float) * filterWidth * filterWidth;
checkCudaErrors(cudaMalloc(&d_filter, filtersize));

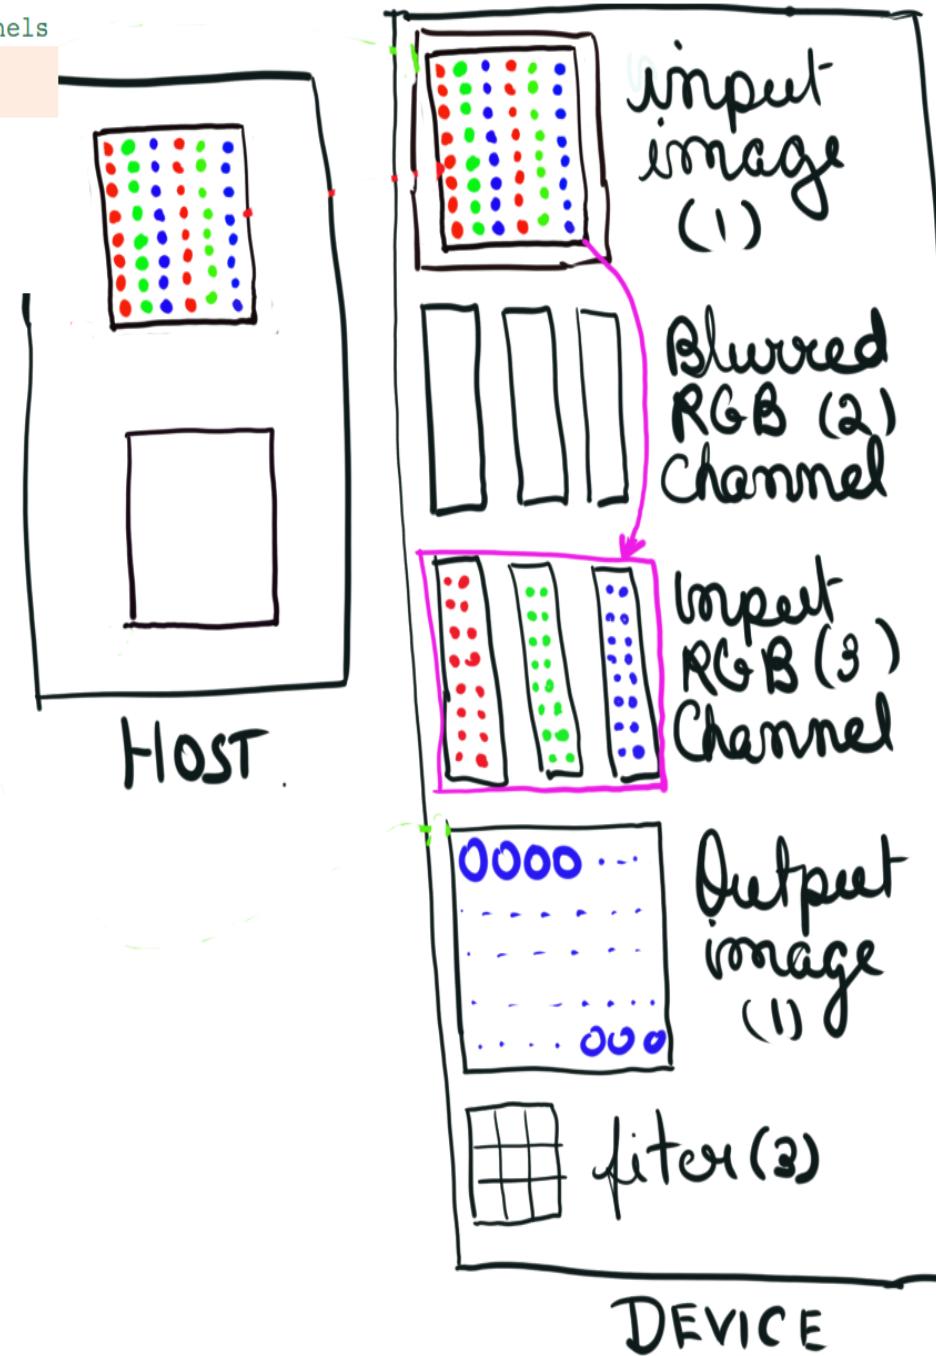
checkCudaErrors(
    cudaMemcpy(d_filter, h_filter, filtersize, cudaMemcpyHostToDevice));
```



CUDA:GPU:MultiDim:ex-3(Blur)

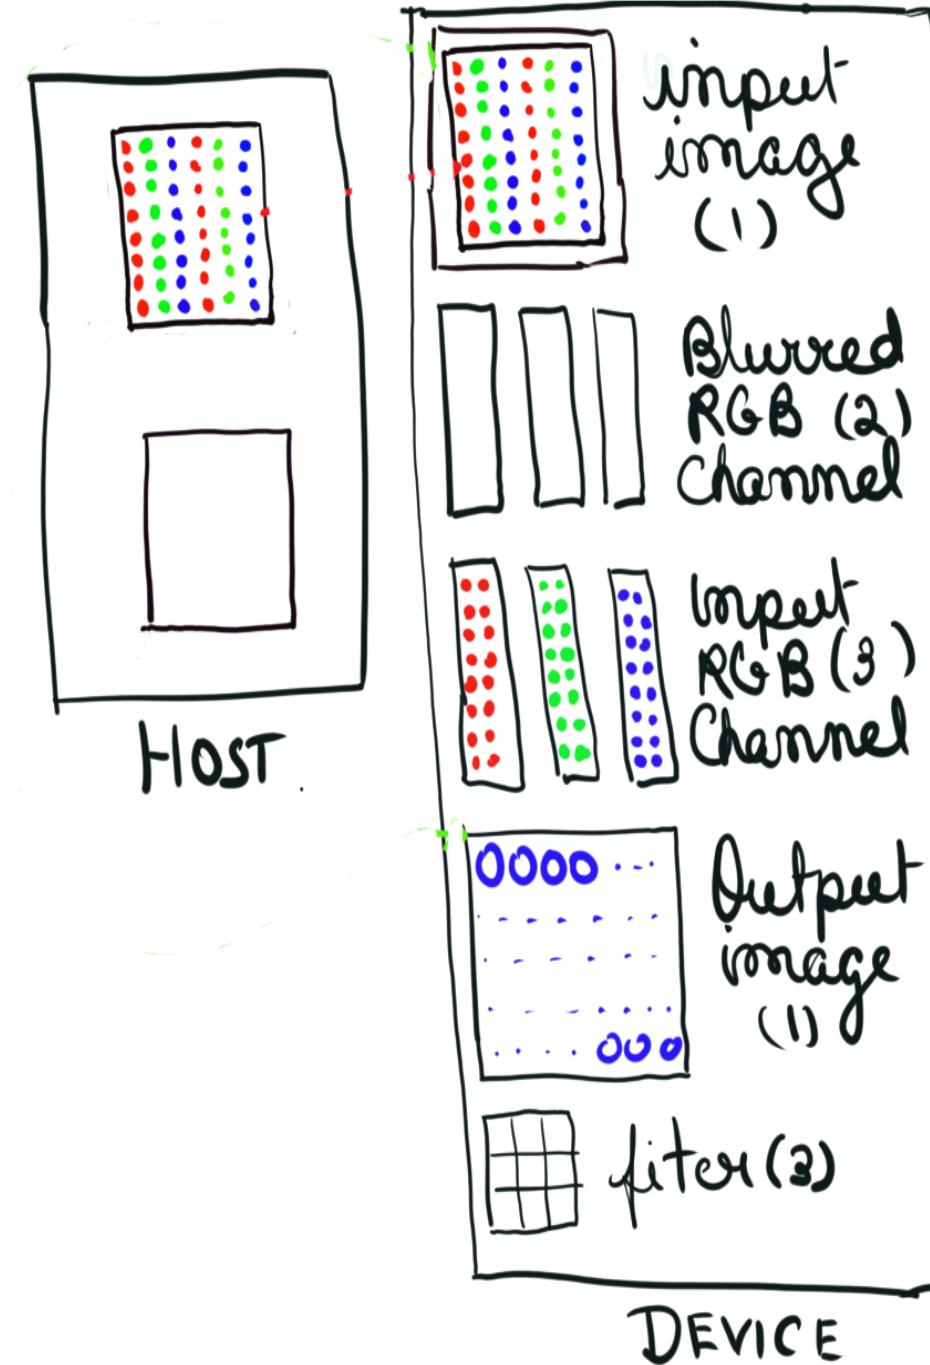
```
const dim3 blockSize(16, 16, 1); //TODO  
const dim3 gridSize(ceil(numCols / 16.0), ceil(numRows / 16.0), 1); //TODO  
//TODO: Launch a kernel for separating the RGBA image into different color channels  
separateChannels<<<gridSize, blockSize>>>(d_inputImageRGBA, numRows,  
    numCols, d_red, d_green, d_blue);  
  
// Call cudaDeviceSynchronize(), then call checkCudaErrors() immediately after  
// launching your kernel to make sure that you didn't make any mistakes.  
cudaDeviceSynchronize();  
checkCudaErrors(cudaGetLastError());
```

```
_global_  
void separateChannels(const uchar4* const inputImageRGBA, int numRows,  
    int numCols, unsigned char* const redChannel,  
    unsigned char* const greenChannel, unsigned char* const blueChannel) {  
  
    int Col = threadIdx.x + blockIdx.x * blockDim.x;  
    int Row = threadIdx.y + blockIdx.y * blockDim.y;  
    int i = Row * numCols + Col;  
    if (Col < numCols && Row < numRows) {  
        uchar4 rgba = inputImageRGBA[i];  
        redChannel[i] = rgba.x;  
        greenChannel[i] = rgba.y;  
        blueChannel[i] = rgba.z;  
    }  
}
```



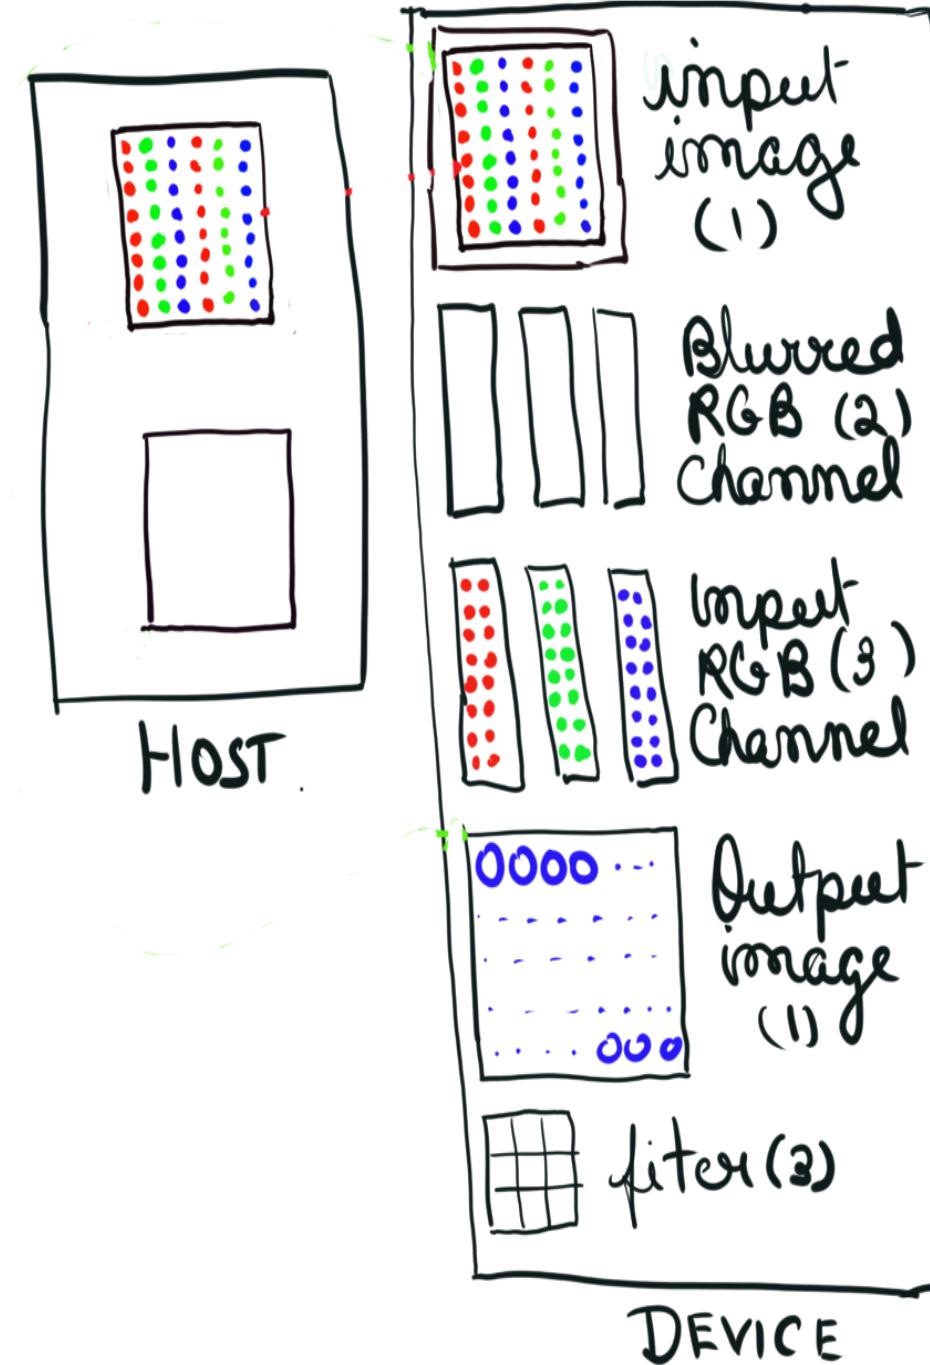
CUDA:GPU:MultiDim:ex-3(Blur)

```
gaussian_blur<<<gridSize, blockSize>>>(d_red, d_redBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
gaussian_blur<<<gridSize, blockSize>>>(d_green, d_greenBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
gaussian_blur<<<gridSize, blockSize>>>(d_blue, d_blueBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
global  
void gaussian blur(const unsigned char* const inputChannel,  
    unsigned char* const outputChannel, int numRows, int numCols,  
    const float* const filter, const int filterWidth) {  
    int col = blockIdx.x * blockDim.x + threadIdx.x;  
    int row = blockIdx.y * blockDim.y + threadIdx.y;  
    if (col >= numCols || row >= numRows) {  
        return;  
    }  
    float result = 0.0f;  
    for (int row_delta = -filterWidth / 2; row_delta <= filterWidth / 2;  
        row_delta++) {  
        for (int col_delta = -filterWidth / 2; col_delta <= filterWidth / 2;  
            col_delta++) {  
            // Compute the coordinates of the value this coefficient applies  
            // Apply clamping to image boundaries.  
            int value_row = min(max(row + row_delta, 0), numRows - 1);  
            int value_col = min(max(col + col_delta, 0), numCols - 1);  
  
            // Compute the partial sum this value adds to the result when sc  
            // the appropriate coefficient.  
            float channel_value = static_cast<float>(inputChannel[value_row  
                * numCols + value_col]);  
            float filter_coefficient = filter[(row_delta + filterWidth / 2)  
                * filterWidth + (col_delta + filterWidth / 2)];  
            result += channel_value * filter_coefficient;  
        }  
    }  
    outputChannel[row * numCols + col] = result;  
}
```



CUDA:GPU:MultiDim:ex-3(Blur)

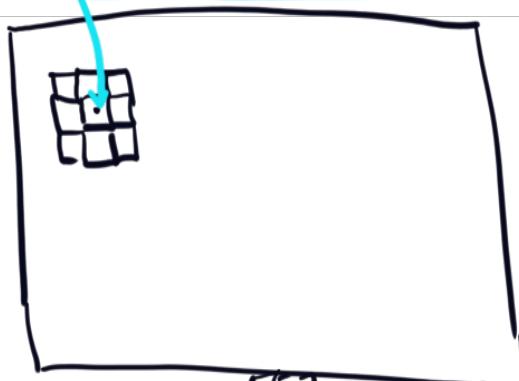
```
gaussian_blur<<<gridSize, blockSize>>>(d_red, d_redBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
gaussian_blur<<<gridSize, blockSize>>>(d_green, d_greenBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
gaussian_blur<<<gridSize, blockSize>>>(d_blue, d_blueBlurred, numRows,  
    numCols, d_filter, filterWidth);  
  
global  
void gaussian blur(const unsigned char* const inputChannel,  
    unsigned char* const outputChannel, int numRows, int numCols,  
    const float* const filter, const int filterWidth) {  
    int col = blockIdx.x * blockDim.x + threadIdx.x;  
    int row = blockIdx.y * blockDim.y + threadIdx.y;  
    if (col >= numCols || row >= numRows) {  
        return;  
    }  
    float result = 0.0f;  
    for (int row_delta = -filterWidth / 2; row_delta <= filterWidth / 2;  
        row_delta++) {  
        for (int col_delta = -filterWidth / 2; col_delta <= filterWidth / 2;  
            col_delta++) {  
            // Compute the coordinates of the value this coefficient applies  
            // Apply clamping to image boundaries.  
            int value_row = min(max(row + row_delta, 0), numRows - 1);  
            int value_col = min(max(col + col_delta, 0), numCols - 1);  
  
            // Compute the partial sum this value adds to the result when sc  
            // the appropriate coefficient.  
            float channel_value = static_cast<float>(inputChannel[value_row  
                * numCols + value_col]);  
            float filter_coefficient = filter[(row_delta + filterWidth / 2)  
                * filterWidth + (col_delta + filterWidth / 2)];  
            result += channel_value * filter_coefficient;  
        }  
    }  
    outputChannel[row * numCols + col] = result;  
}
```



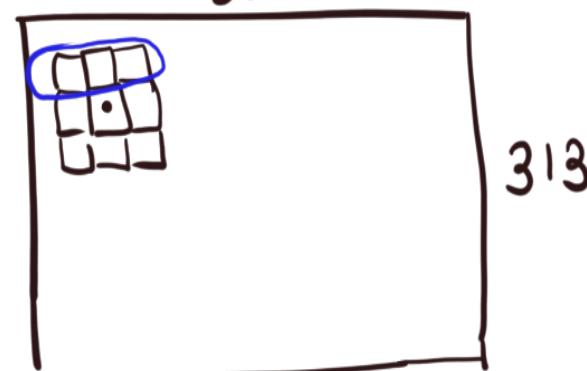
CUDA:GPU:MultiDim:ex-3(Blur)

```
_global_
void gaussian_blur(const unsigned char* const inputChannel,
    unsigned char* const outputChannel, int numRows, int numCols,
    const float* const filter, const int filterWidth) {
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    if (col >= numCols || row >= numRows) {
        return;
    }
    float result = 0.0f;
    for (int row_delta = -filterWidth / 2; row_delta <= filterWidth / 2;
        row_delta++) {
        for (int col_delta = -filterWidth / 2; col_delta <= filterWidth / 2;
            col_delta++) {
            // Compute the coordinates of the value this coefficient applies
            // Apply clamping to image boundaries.
            int value_row = min(max(row + row_delta, 0), numRows - 1);
            int value_col = min(max(col + col_delta, 0), numCols - 1);

            // Compute the partial sum this value adds to the result when sc
            // the appropriate coefficient.
            float channel_value = static_cast<float>(inputChannel[value_row
                * numCols + value_col]);
            float filter_coefficient = filter[(row_delta + filterWidth / 2)
                * filterWidth + (col_delta + filterWidth / 2)];
            result += channel_value * filter_coefficient;
        }
    }
    outputChannel[row * numCols + col] = result;
}
```



- for threadIdx.x = 3 and threadIdx.y = 3 and picture of 557×313 dim.



- channel value index is going to be $(557+557+3-1) = (1116, 1117, 1118, 1673, 1674, 1675, \dots)$
- Filter coefficient index is going to be 0 1 2 4 5 6 7 8
- After adding up writing the result at the same offset.

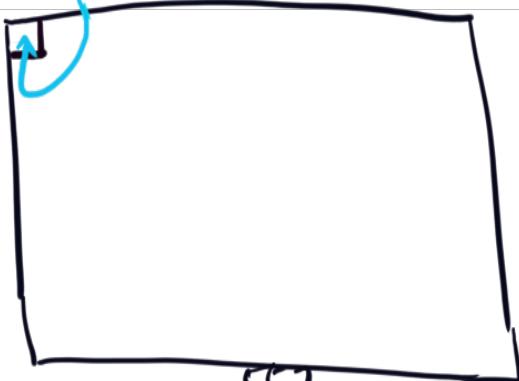
CUDA:GPU:MultiDim:ex-3(Blur)

```
global_
void gaussian_blur(const unsigned char* const inputChannel,
    unsigned char* const outputChannel, int numRows, int numCols,
    const float* const filter, const int filterWidth) {
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    if (col >= numCols || row >= numRows) {
        return;
    }
    float result = 0.0f;
    for (int row_delta = -filterWidth / 2; row_delta <= filterWidth / 2;
        row_delta++) {
        for (int col_delta = -filterWidth / 2; col_delta <= filterWidth / 2;
            col_delta++) {
            // Compute the coordinates of the value this coefficient applies
            // to.
            int valueRow = min(max(row + row_delta, 0), numRows - 1);
            int valueCol = min(max(col + col_delta, 0), numCols - 1);

            // Compute the partial sum this value adds to the result when scaled
            // by the appropriate coefficient.
            float channelValue = static_cast<float>(inputChannel[valueRow *
                numCols + valueCol]);
            float filterCoefficient = filter[(row_delta + filterWidth / 2) *
                filterWidth + (col_delta + filterWidth / 2)];
            result += channelValue * filterCoefficient;
        }
    }
    outputChannel[row * numCols + col] = result;
}
```

3 for
the
picture

- for threadIdx.x = 0 and threadIdx.y = 0 and picture of 557×313 dim.

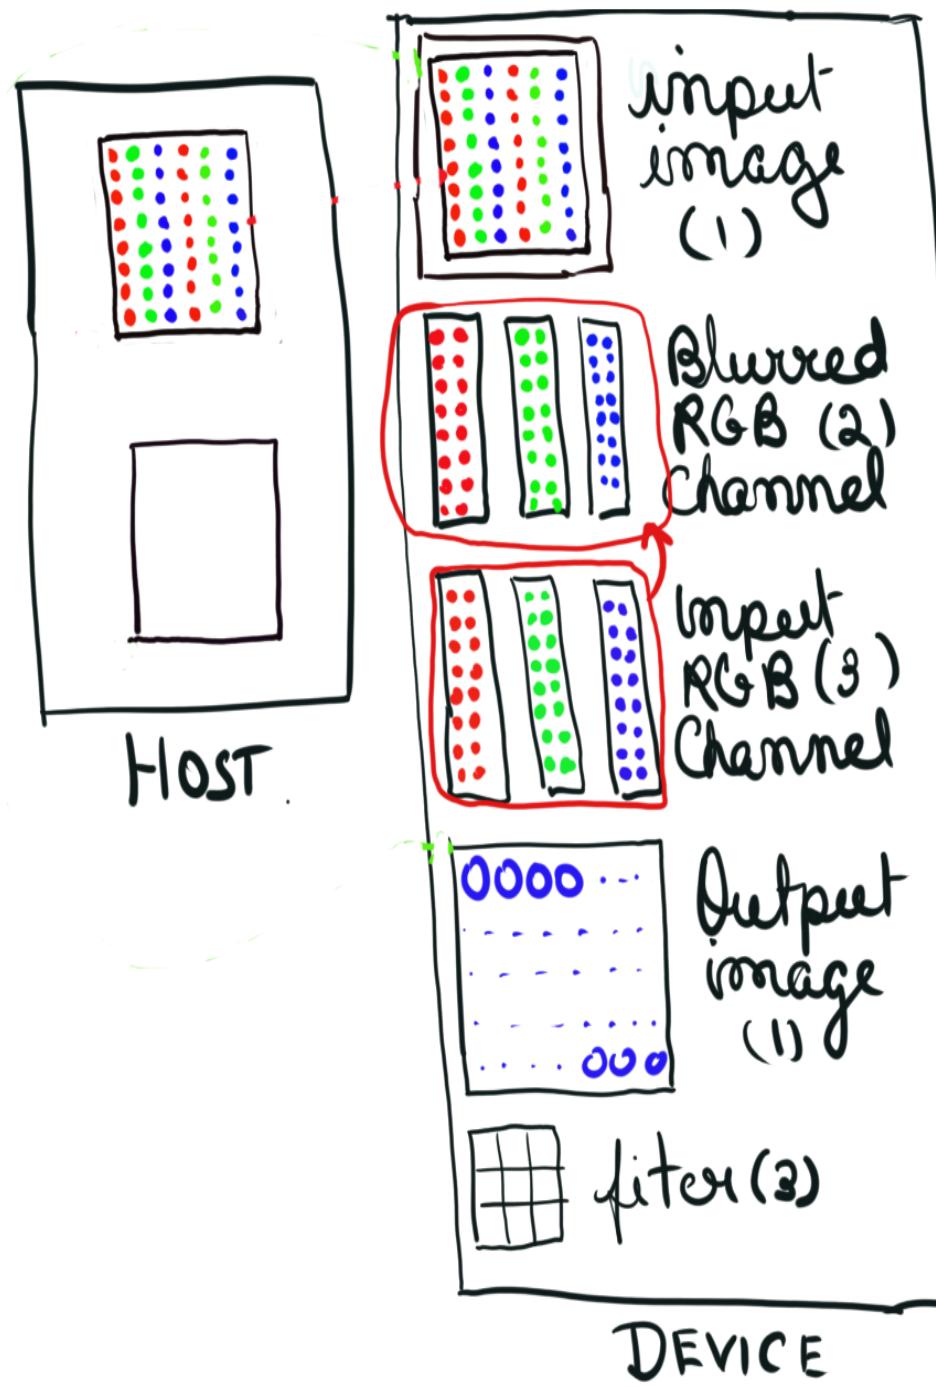


- channel value index is going to be $(0, 0, 1)$, $(0, 1, 1)$, $(557, 557, 558)$.
- Pink values are copied because these are edge pixels

CUDA:GPU:MultiDim:ex-3(Blur)

```
global_
void gaussian_blur(const unsigned char* const inputChannel,
                    unsigned char* const outputChannel, int numRows, int numCols,
                    const float* const filter, const int filterWidth) {
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    if (col >= numCols || row >= numRows) {
        return;
    }
    float result = 0.0f;
    for (int row_delta = -filterWidth / 2; row_delta <= filterWidth / 2;
         row_delta++) {
        for (int col_delta = -filterWidth / 2; col_delta <= filterWidth / 2;
             col_delta++) {
            // Compute the coordinates of the value this coefficient applies
            // Apply clamping to image boundaries.
            int value_row = min(max(row + row_delta, 0), numRows - 1);
            int value_col = min(max(col + col_delta, 0), numCols - 1);

            // Compute the partial sum this value adds to the result when sc
            // the appropriate coefficient.
            float channel_value = static_cast<float>(inputChannel[value_row
                * numCols + value_col]);
            float filter_coefficient = filter[(row_delta + filterWidth / 2)
                * filterWidth + (col_delta + filterWidth / 2)];
            result += channel_value * filter_coefficient;
        }
    }
    outputChannel[row * numCols + col] = result;
}
```



CUDA:GPU:MultiDim:ex-3(Blur)

```
recombineChannels<<<gridSize, blockSize>>>(d_redBlurred, d_greenBlurred,  
d_blueBlurred, d_outputImageRGBA, numRows, numCols);
```

global

```
void recombineChannels(const unsigned char* const redChannel,  
                      const unsigned char* const greenChannel,  
                      const unsigned char* const blueChannel, uchar4* const outputImageRGBA,  
                      int numRows, int numCols) {  
    const int2 thread_2D_pos = make_int2(blockIdx.x * blockDim.x + threadIdx.x,  
                                         blockIdx.y * blockDim.y + threadIdx.y);  
    const int thread_1D_pos = thread_2D_pos.y * numCols + thread_2D_pos.x;  
    //make sure we don't try and access memory outside the image  
    //by having any threads mapped there return early  
    if (thread_2D_pos.x >= numCols || thread_2D_pos.y >= numRows)  
        return;  
    unsigned char red = redChannel[thread_1D_pos];  
    unsigned char green = greenChannel[thread_1D_pos];  
    unsigned char blue = blueChannel[thread_1D_pos];  
    //Alpha should be 255 for no transparency  
    uchar4 outputPixel = make uchar4(red, green, blue, 255);  
    outputImageRGBA[thread_1D_pos] = outputPixel;  
}
```

7a

} recombine

} The separately blurred RGB pixels.

