

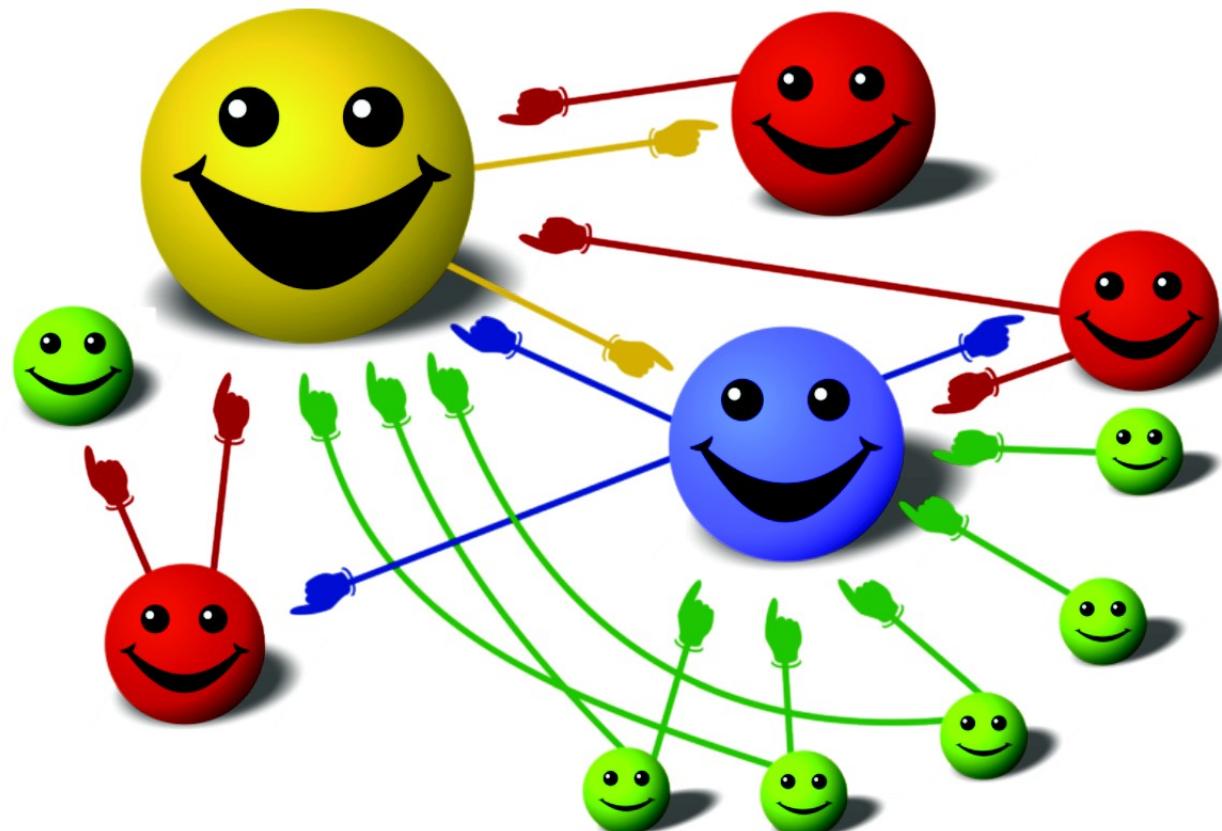
Optimizing Joins-1

By Mohit Kumar

Spark Internals: Pagerank

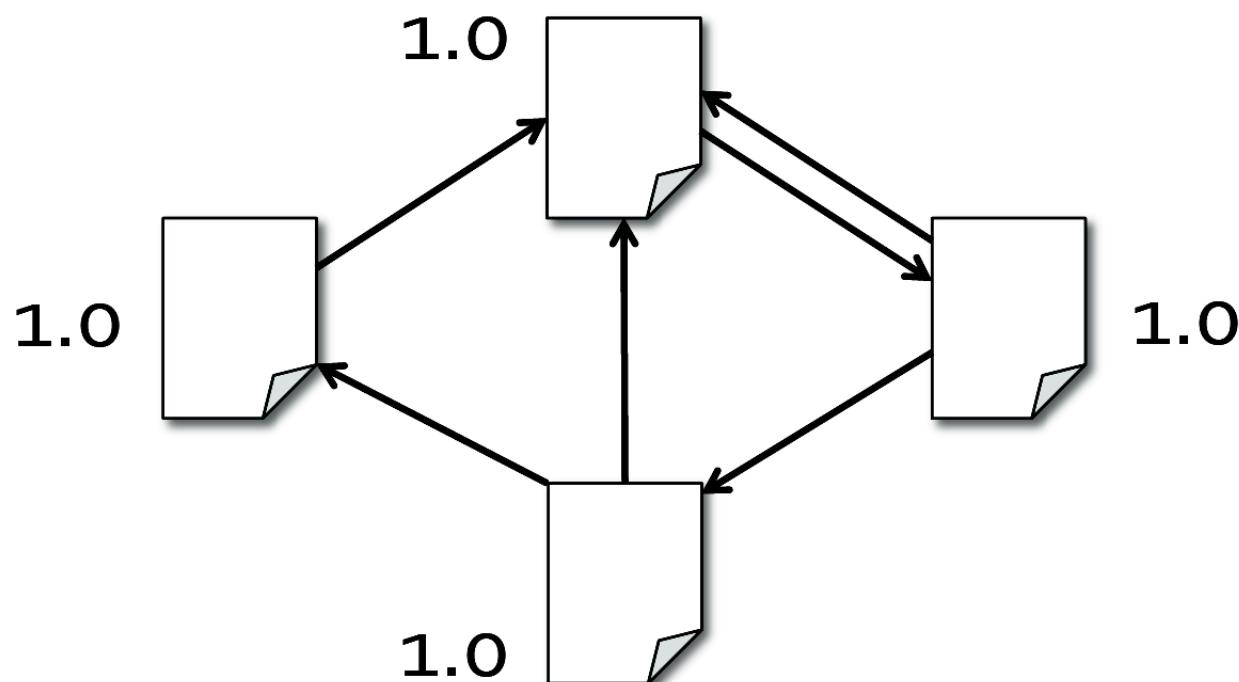
Give pages ranks (scores) based on links to them

- » Links from many pages → high rank
- » Link from a high-rank page → high rank



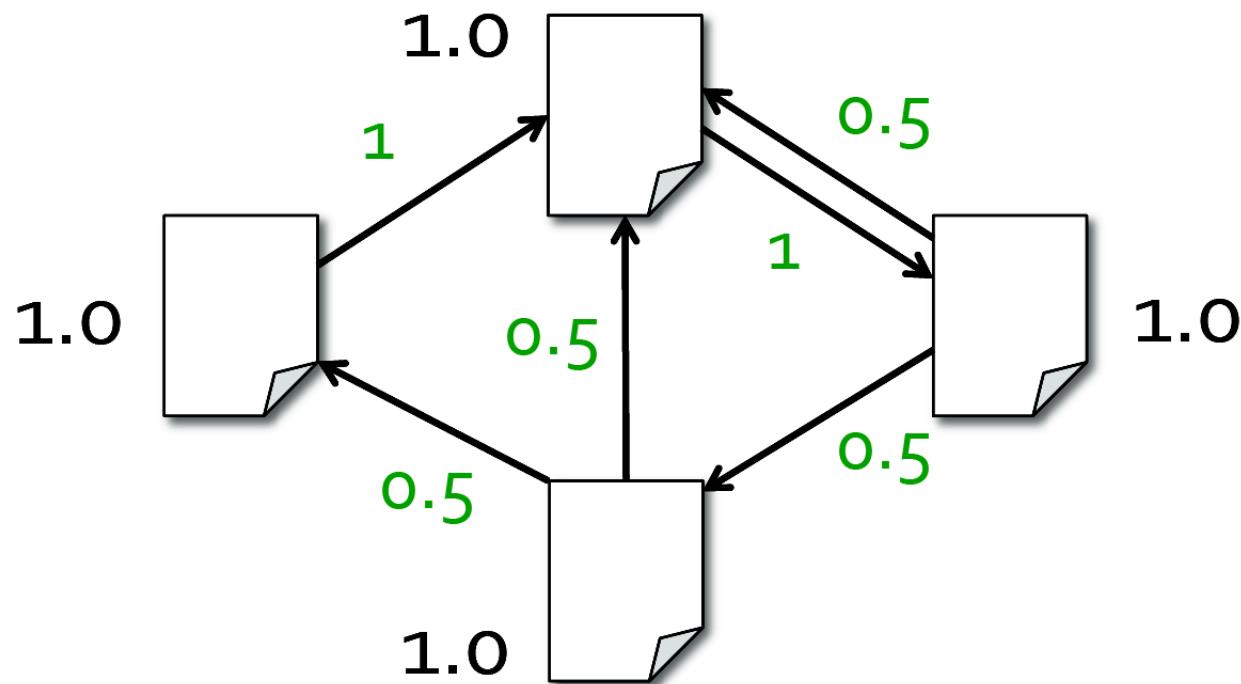
Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contris}$



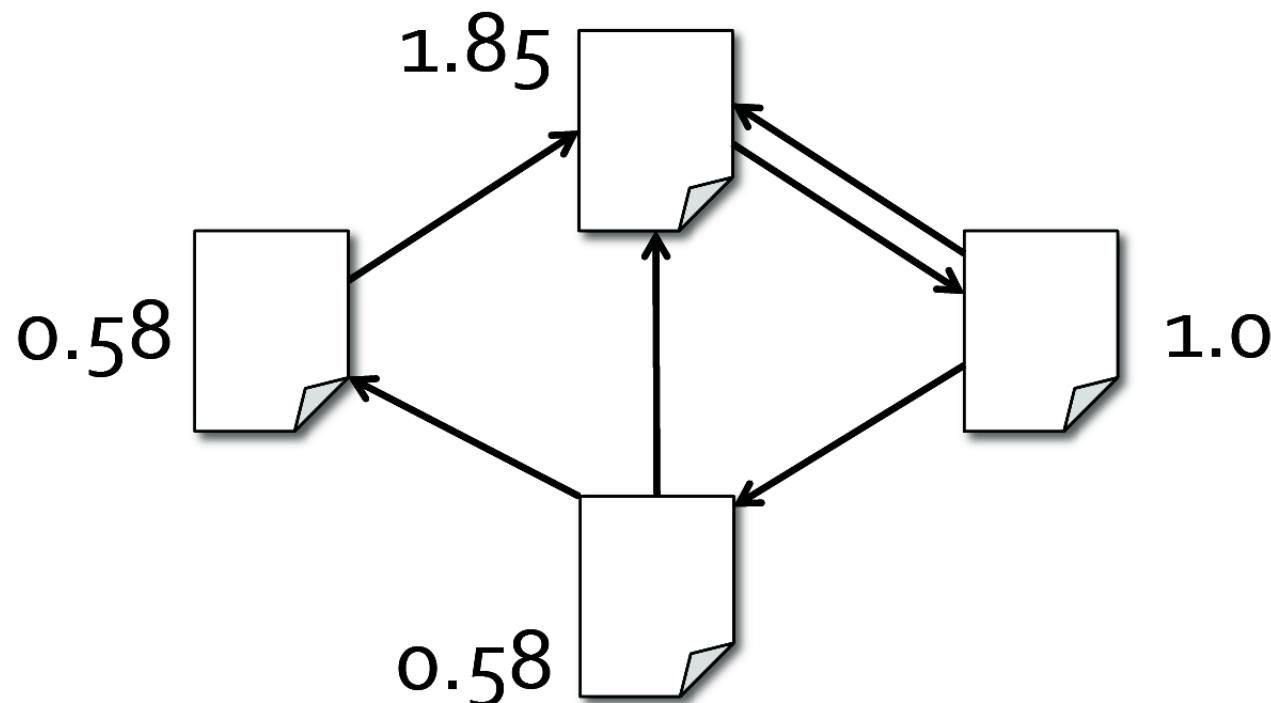
Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



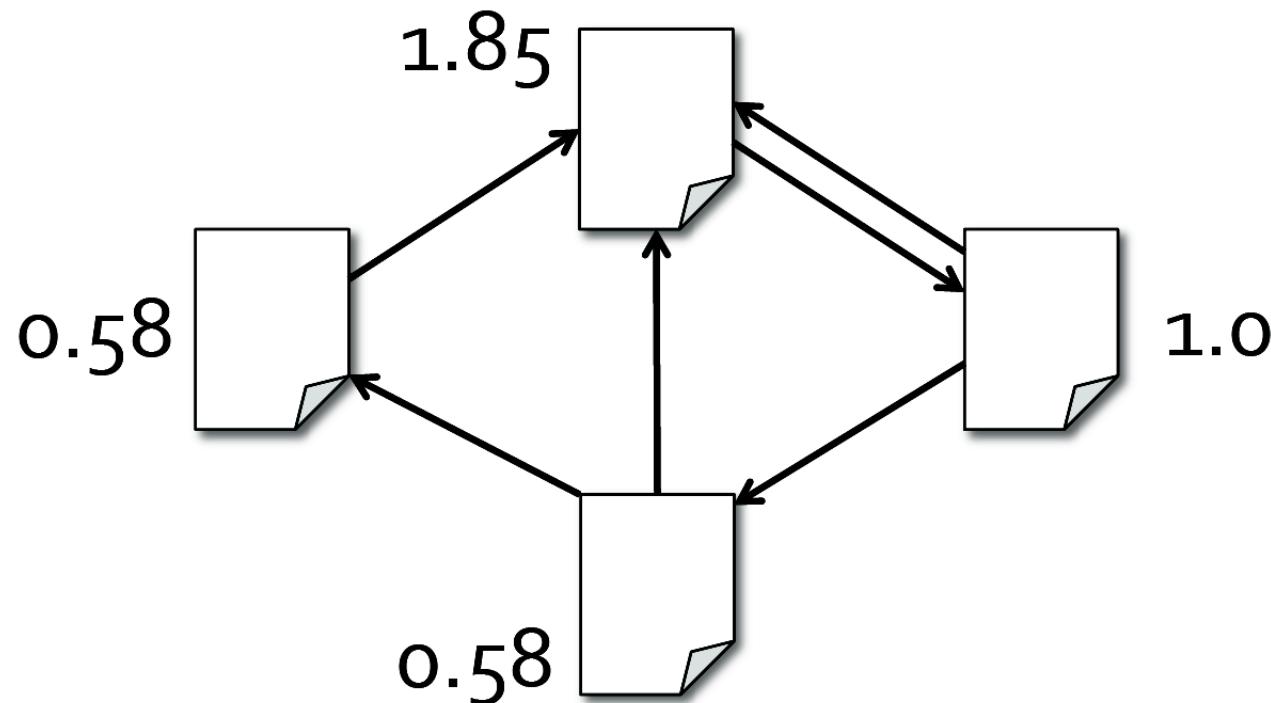
Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



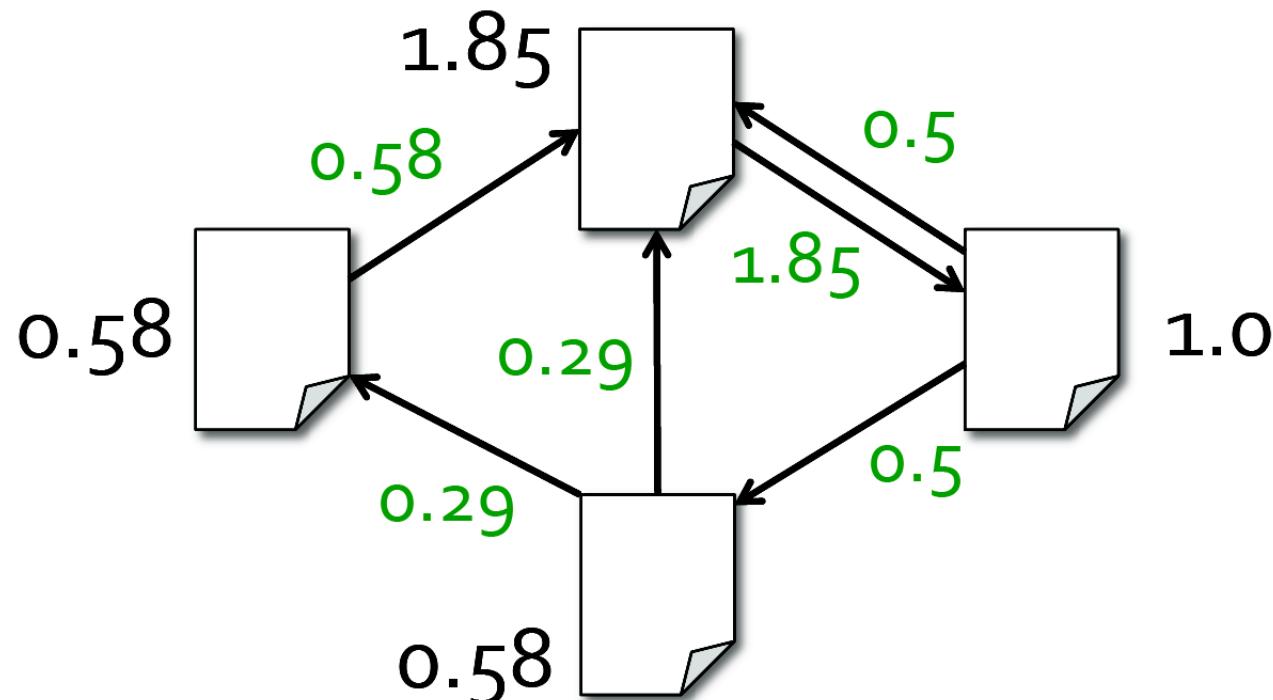
Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



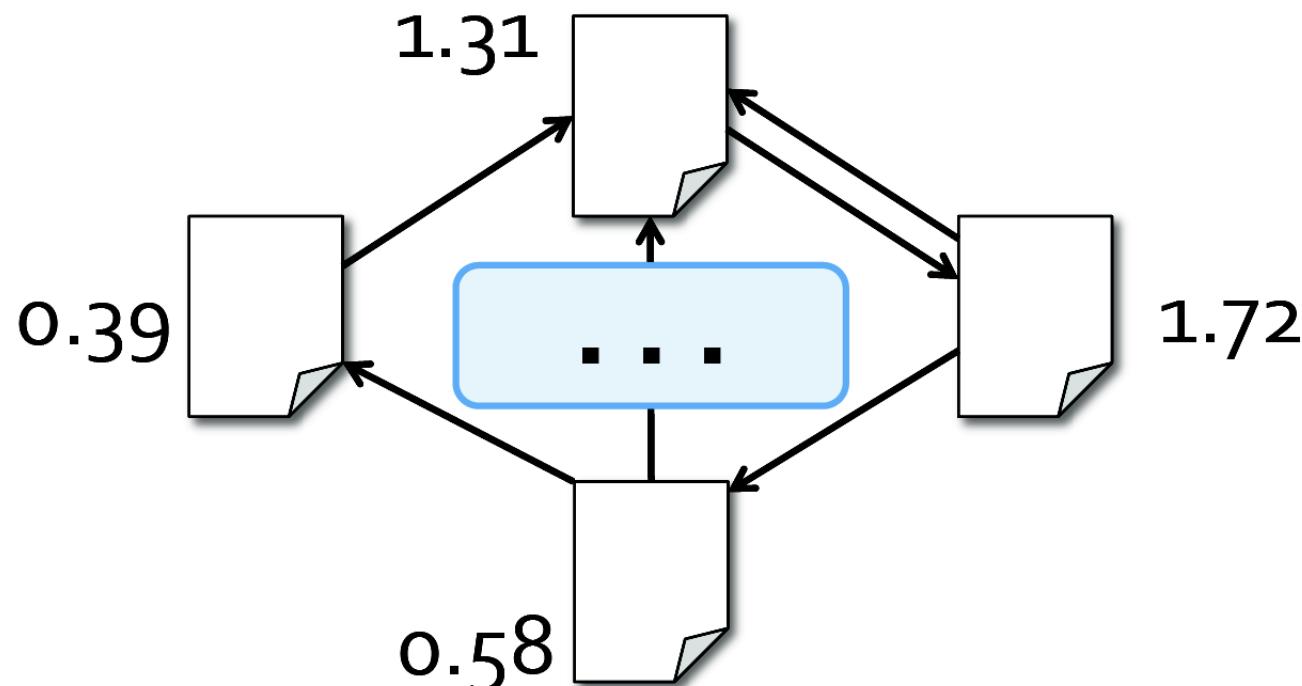
Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Spark Internals: Pagerank

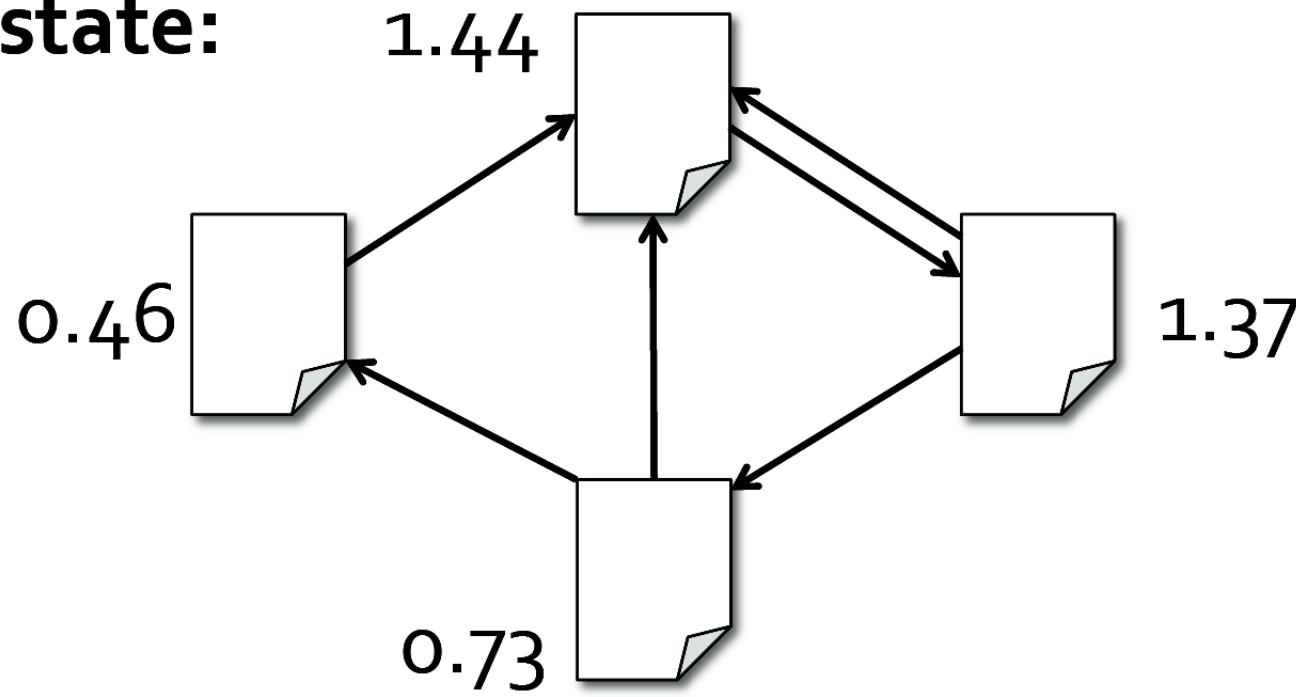
1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$



Spark Internals: Pagerank

1. Start each page at a rank of 1
2. On each iteration, have page p contribute $\text{rank}_p / |\text{neighbors}_p|$ to its neighbors
3. Set each page's rank to $0.15 + 0.85 \times \text{contribs}$

Final state:



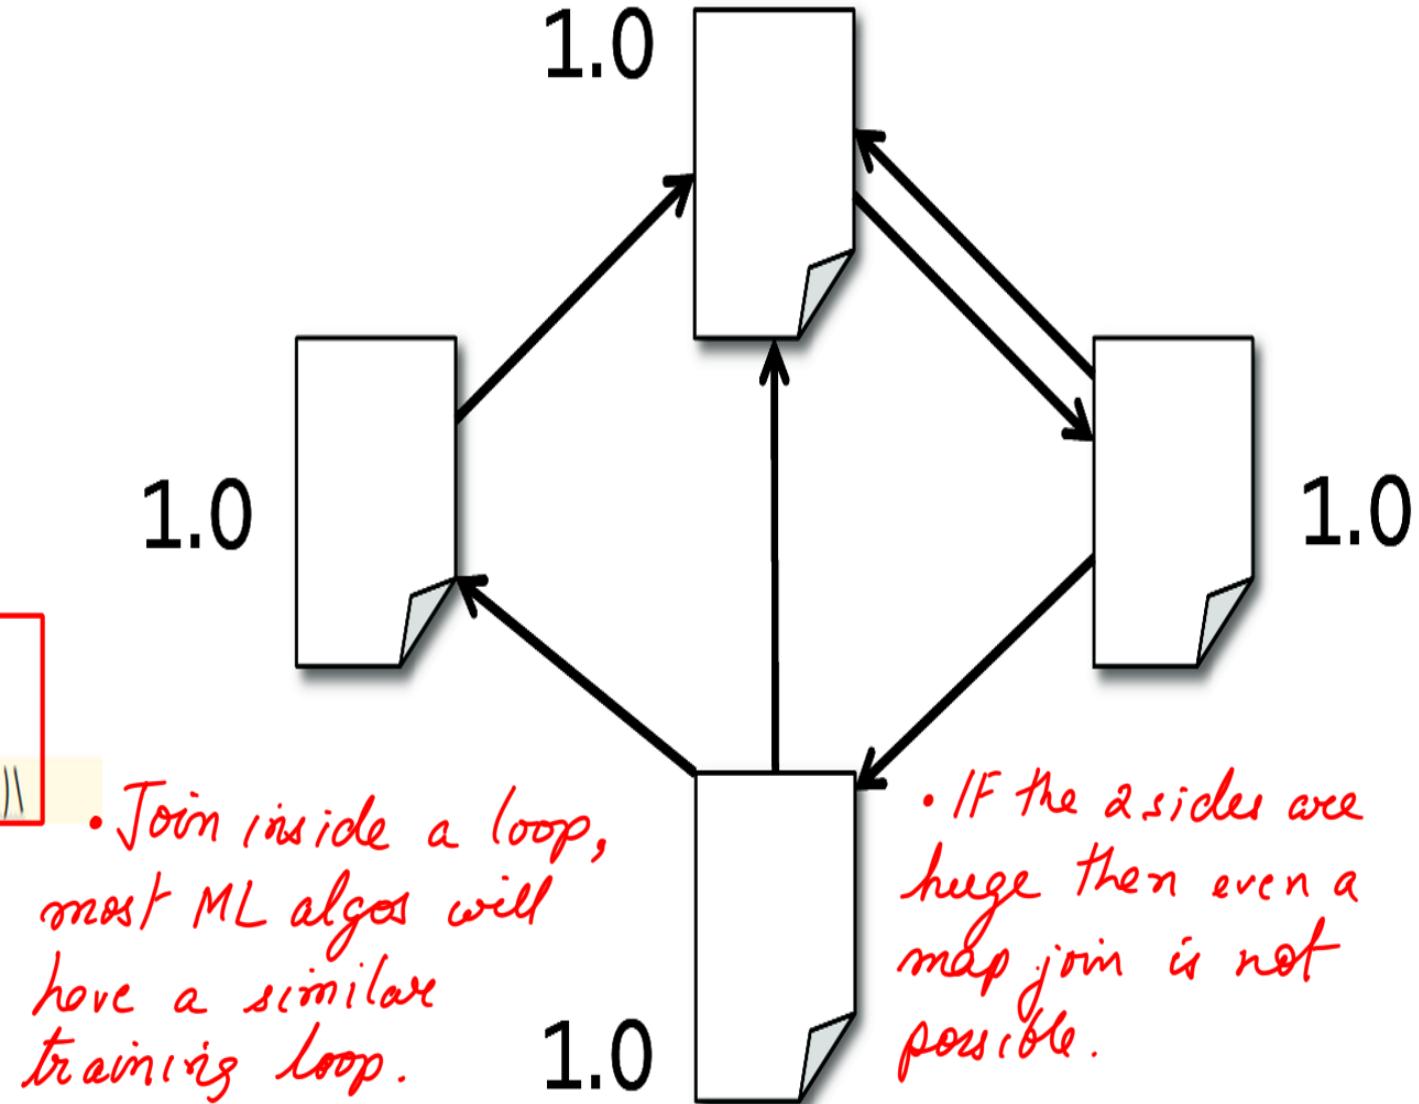
Spark Internals: Pagerank

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

```
ranks=links.map(lambda k: (k[0],1))
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=tolink,(joined[1]/len(joined[0]))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter) :\n    ranks = links.join(ranks).values()\n        .flatMap(lambda joined: processJoined(joined))\n        .reduceByKey(add)\n        .map(lambda x: (x[0],0.15 + x[1]*.85))
```



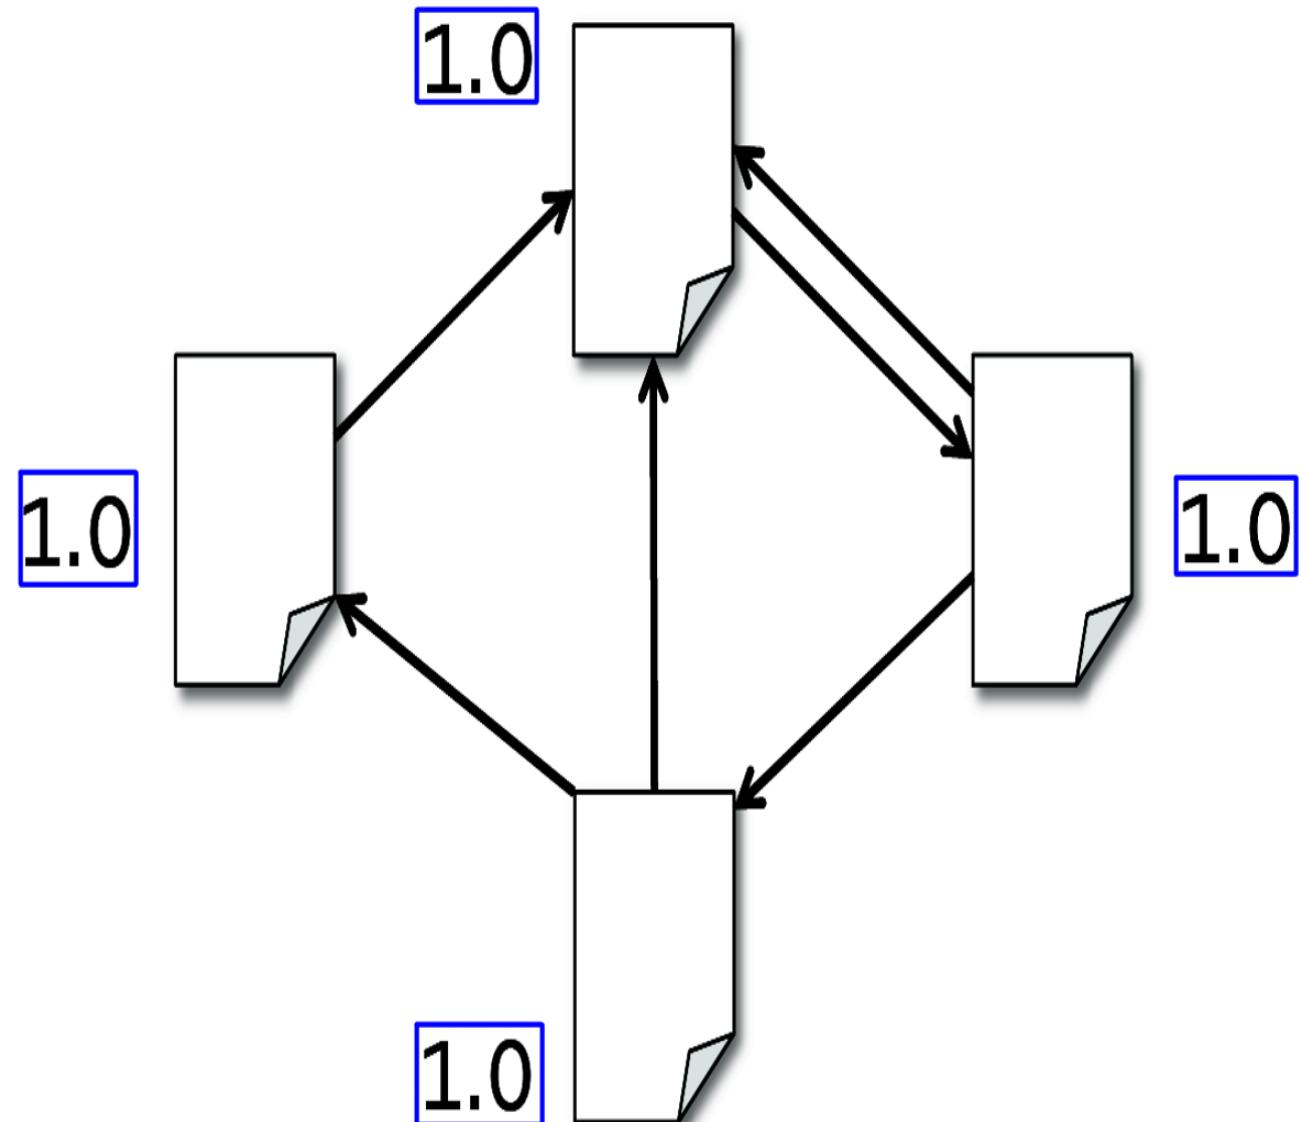
Spark Internals: Pagerank

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

```
ranks=links.map(lambda k: (k[0],1))
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=tolink,(joined[1]/len(joined[0]))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter) :\n    ranks = links.join(ranks).values()\n        .flatMap(lambda joined: processJoined(joined))\n        .reduceByKey(add)\n        .map(lambda x: (x[0],0.15 + x[1]*.85))
```



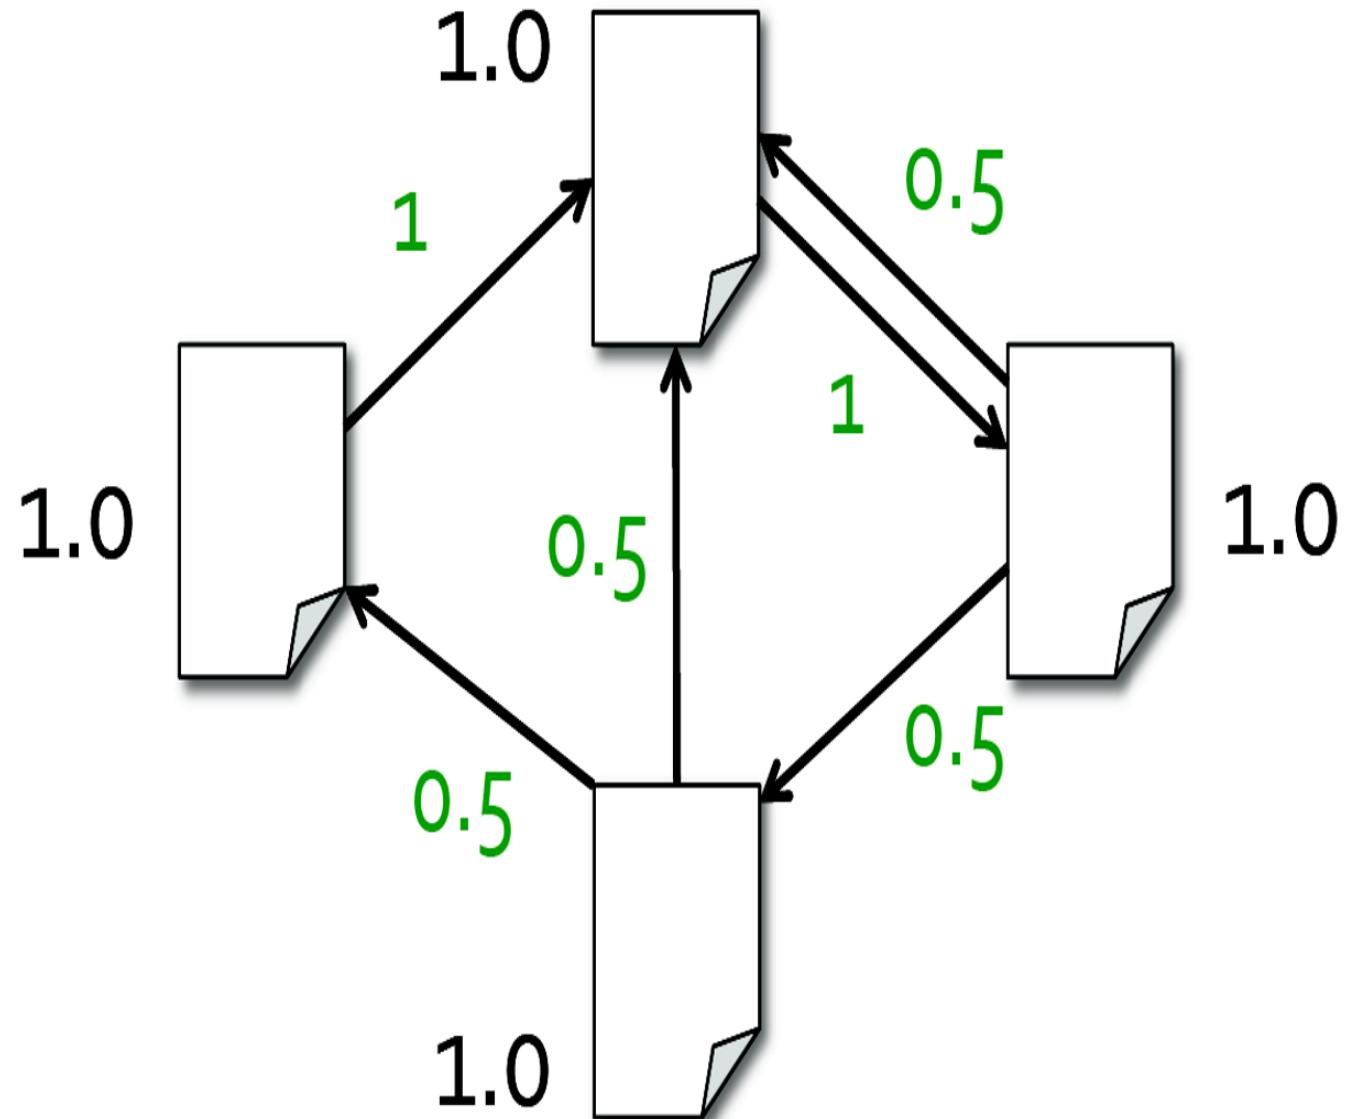
Spark Internals: Pagerank

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

```
ranks=links.map(lambda k: (k[0],1))
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=tolink,(joined[1]/len(joined[0]))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter) :\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .map(lambda x: (x[0],0.15 + x[1]*.85))
```



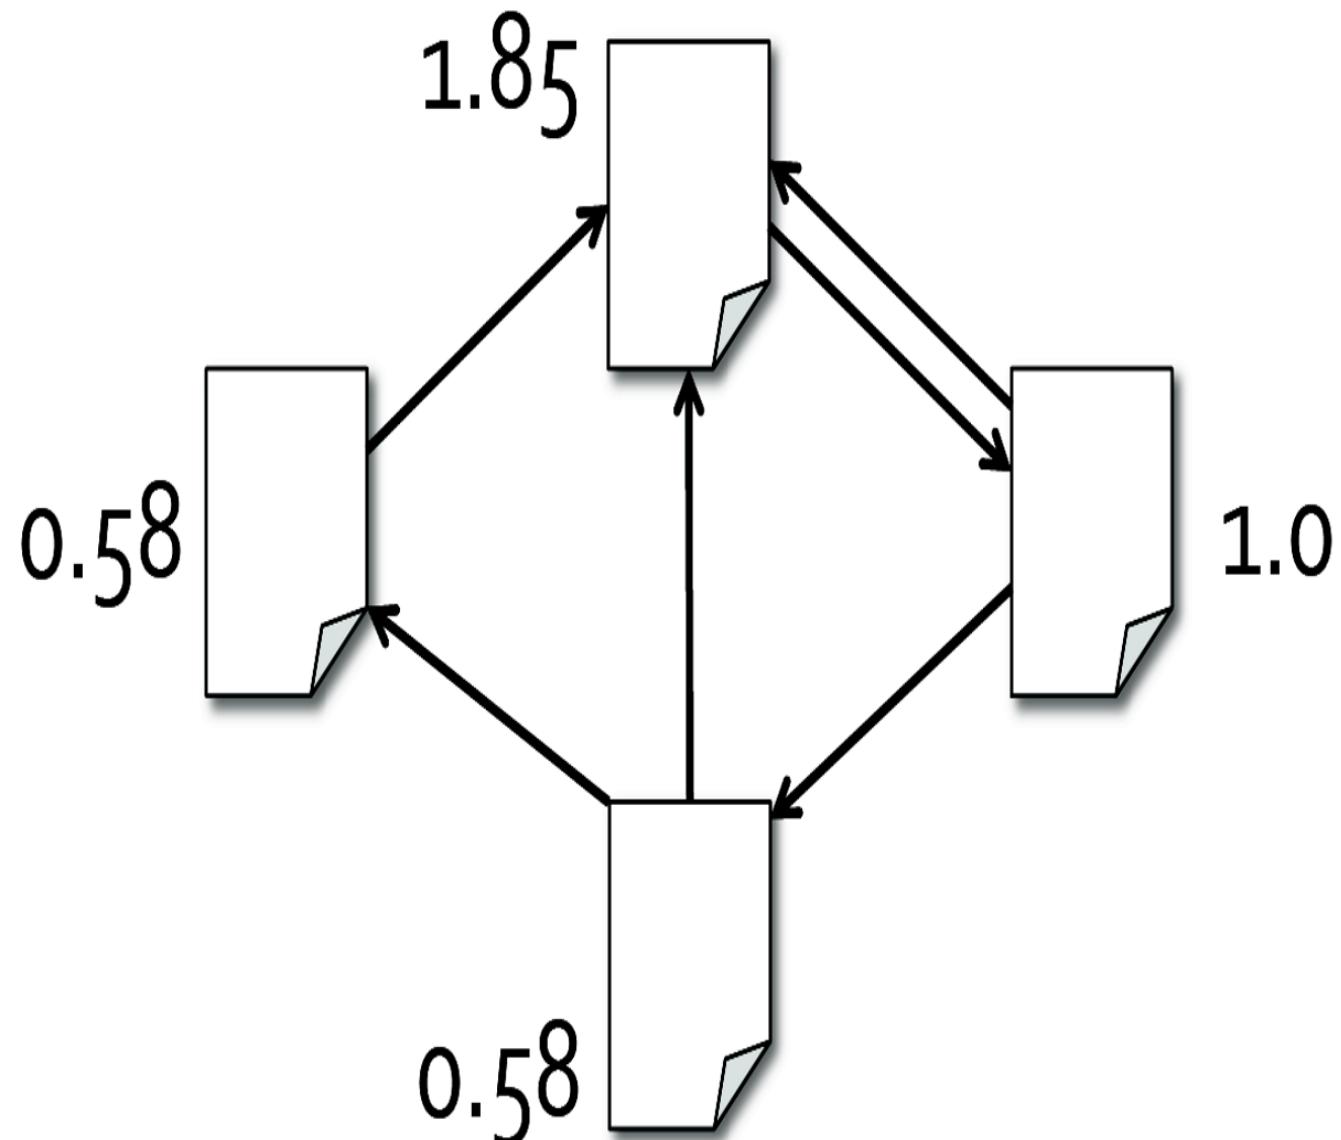
Spark Internals: Pagerank

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

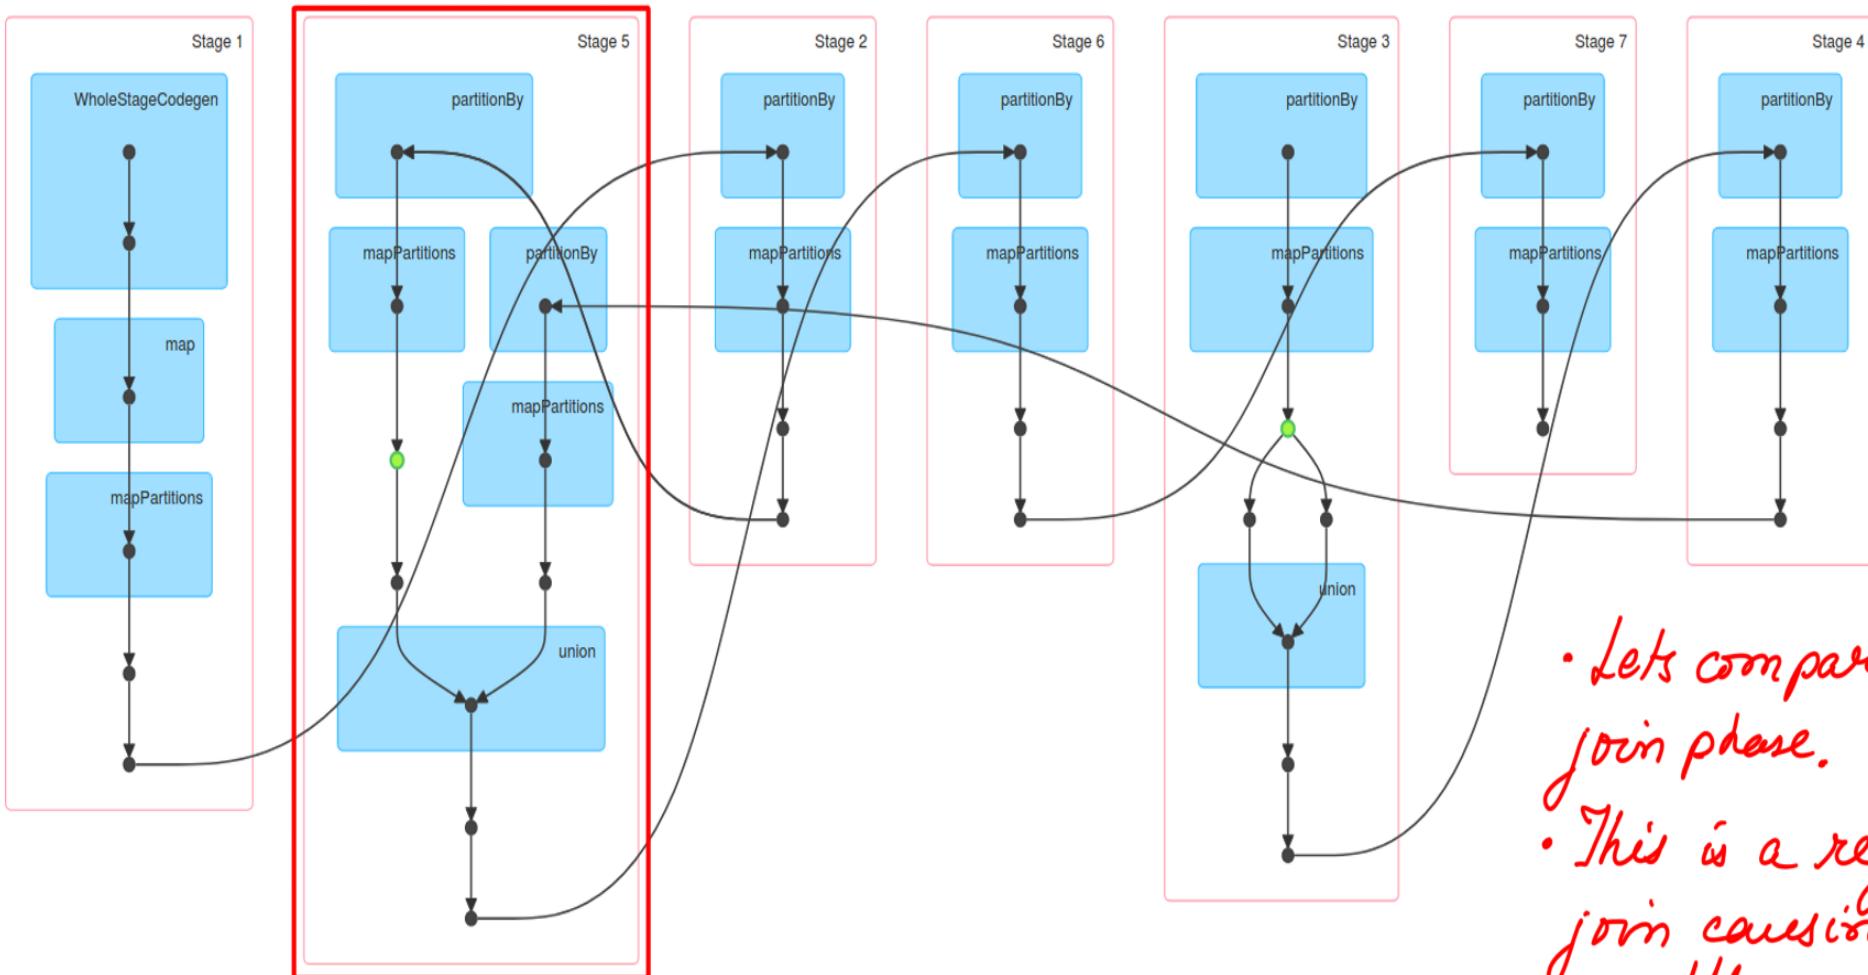
```
ranks=links.map(lambda k: (k[0],1))
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=tolink,(joined[1]/len(joined[0]))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter) :\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .map(lambda x: (x[0],0.15 + x[1]*.85))
```



Spark Internals: Pagerank

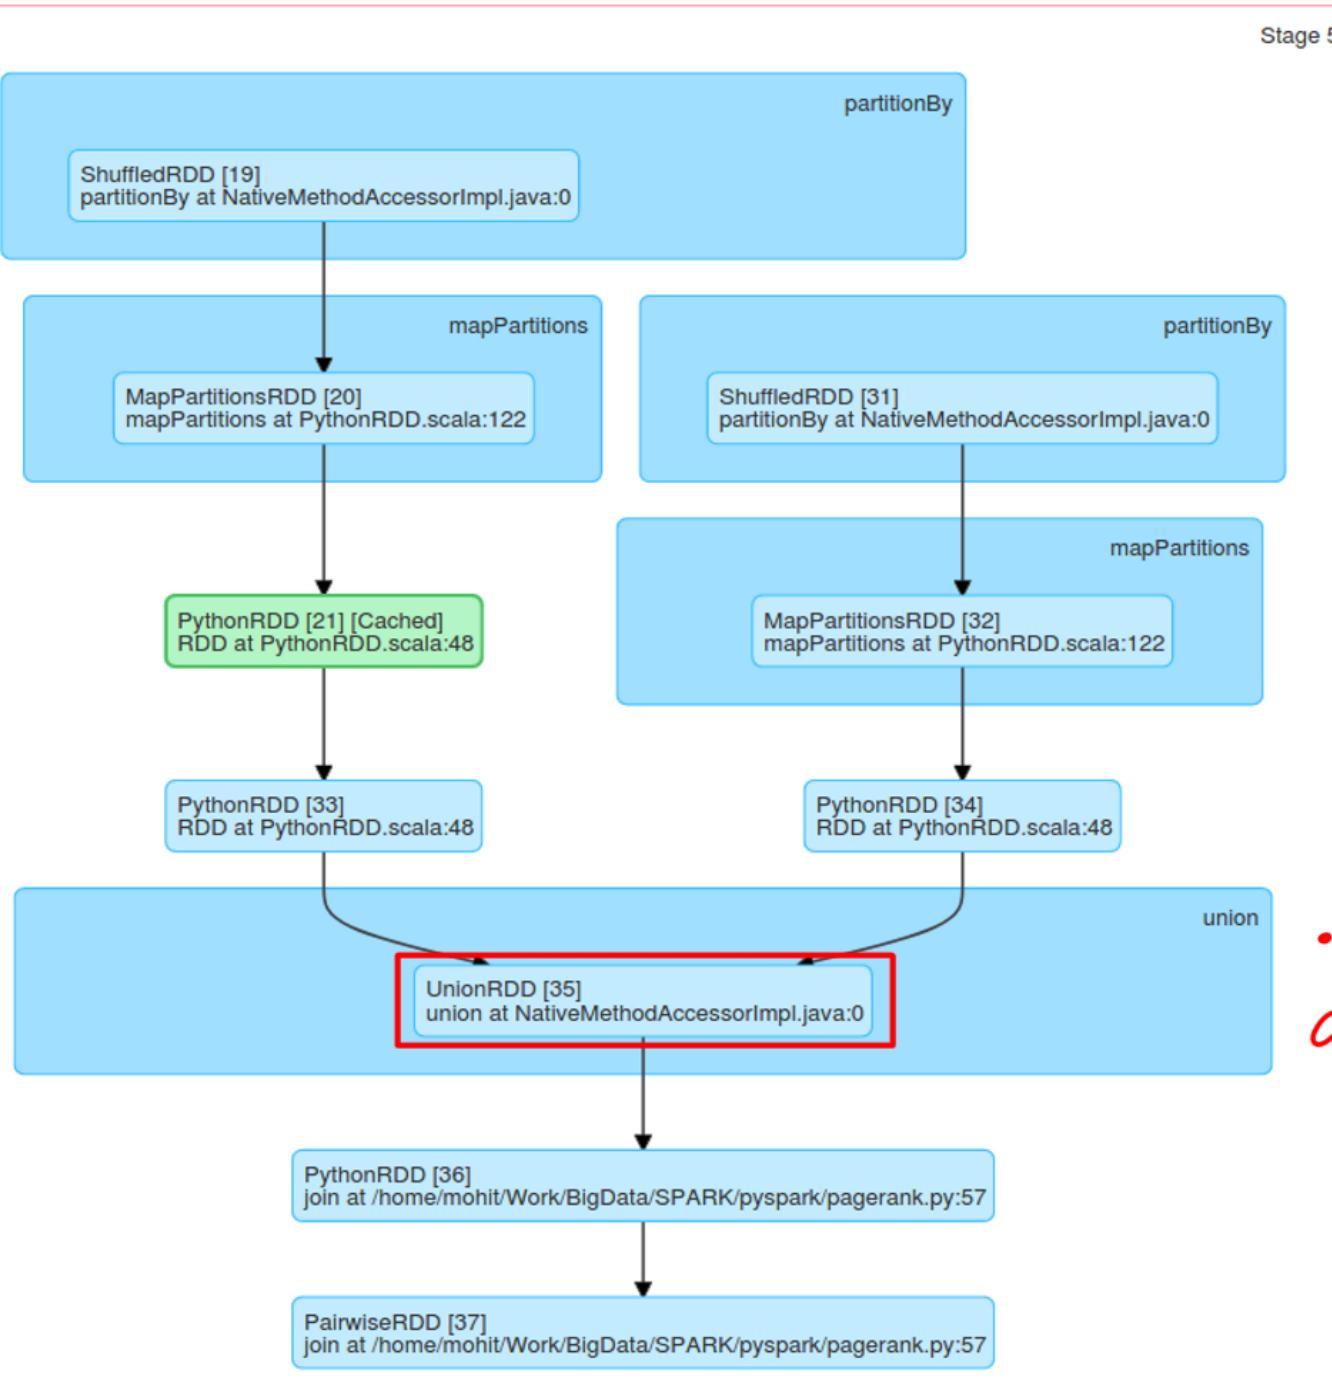


- Lets compare the join phase.
- This is a regular join causing shuffle-sort.

Completed Stages (7)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
7	collect at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:64	+details	2021/03/15 10:15:04	55 ms	3/3		506.0 B	
6	reduceByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:58	+details	2021/03/15 10:15:04	63 ms	3/3		1017.0 B	506.0 B
5	join at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:57	+details	2021/03/15 10:15:04	68 ms	3/3	310.0 B	334.0 B	1017.0 B
4	reduceByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:58	+details	2021/03/15 10:15:04	64 ms	2/2		768.0 B	334.0 B
3	join at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:57	+details	2021/03/15 10:15:04	0.1 s	2/2	310.0 B	185.0 B	768.0 B
2	groupByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:33	+details	2021/03/15 10:15:04	75 ms	1/1		209.0 B	185.0 B
1	distinct at /home/mohit/Work/BigData/SPARK/pyspark/pagerank.py:33	+details	2021/03/15 10:15:03	0.5 s	1/1			209.0 B

Spark Internals: Pagerank



• Regular SMJoin with
a UnionRDD.

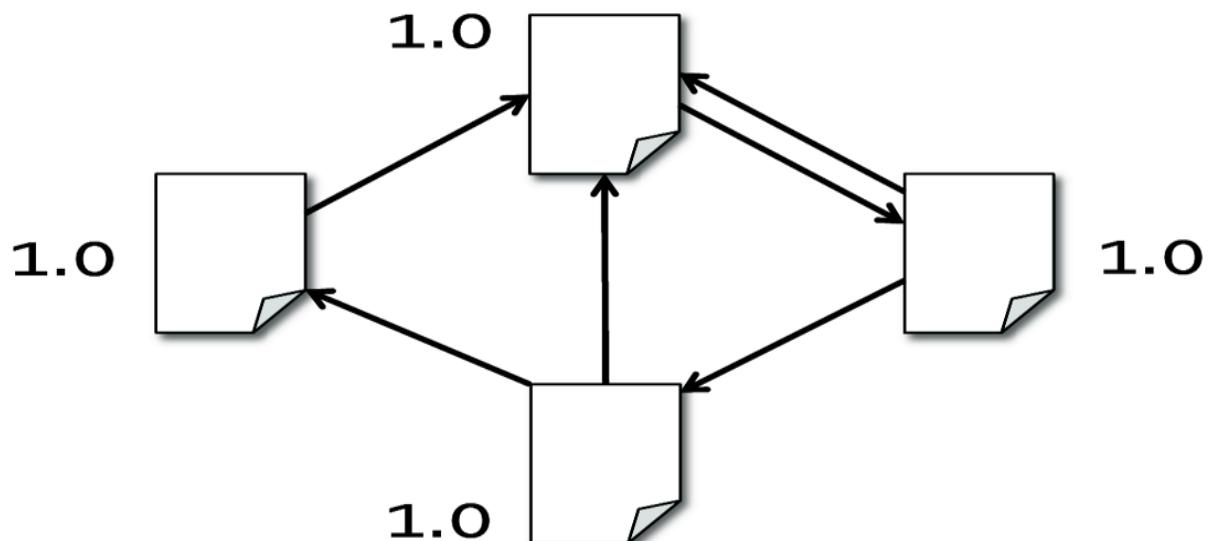
Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=(tolink,(joined[1]/len(joined[0])))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter):\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .mapValues(lambda x: 0.15 + x * 0.85)
```



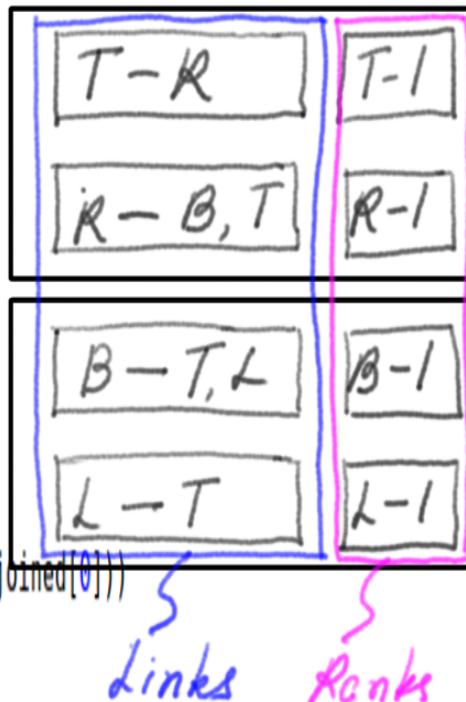
Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n    .distinct()\n    .groupByKey()\n    .cache();
```

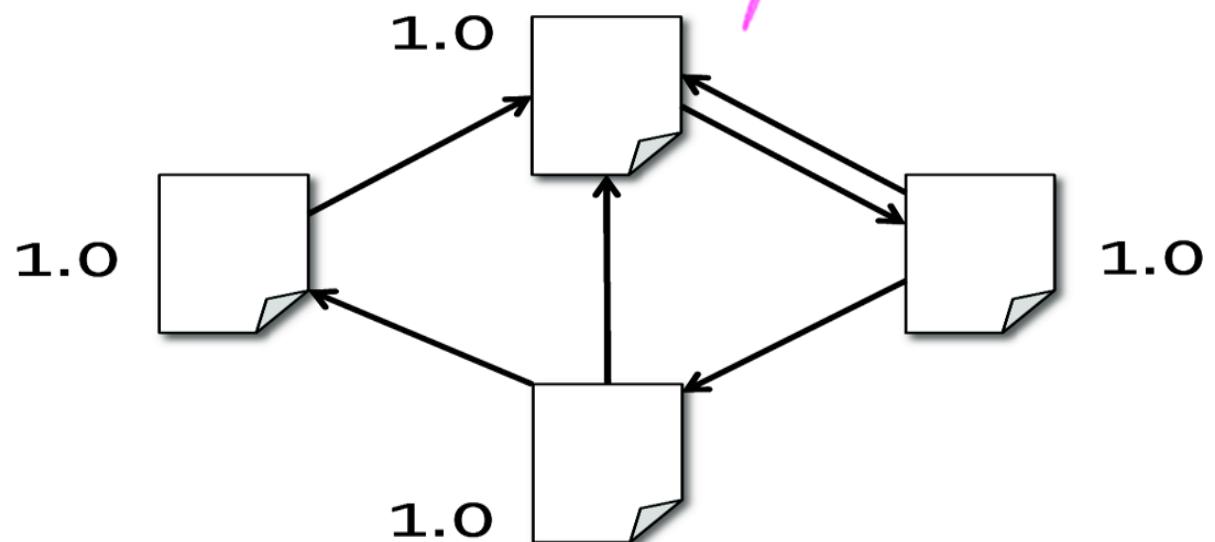
```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=tolink,(joined[1]/len(joined[0]))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter) :\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n        .flatMap(lambda joined: processJoined(joined))\n        .reduceByKey(add)\n        .mapValues(lambda x: 0.15 + x * 0.85)
```



• Do you notice anything special??



Spark Internals: Pagerank: Optimized

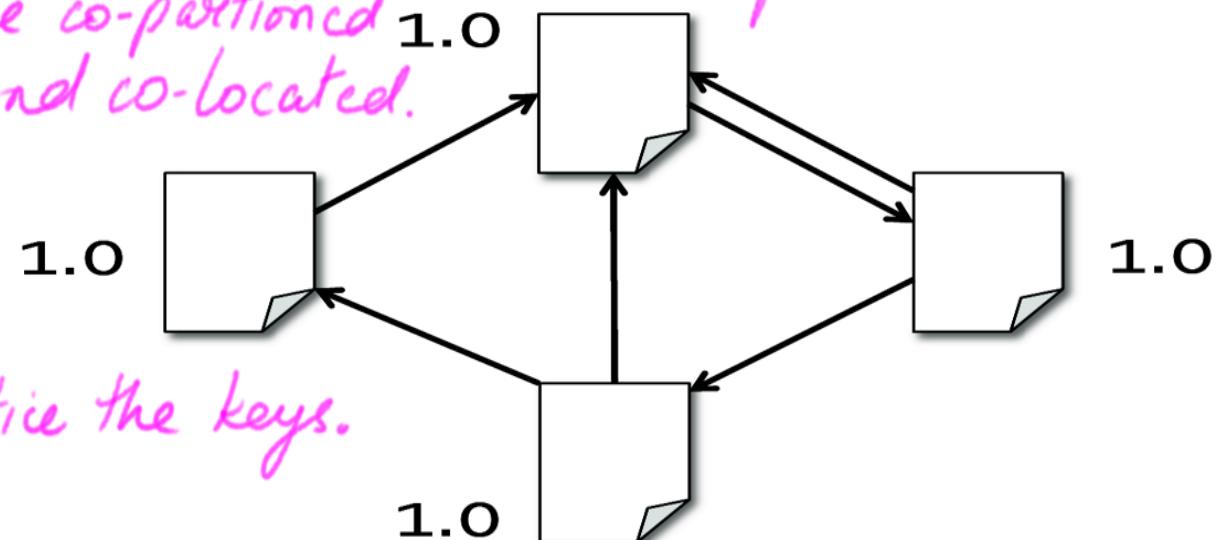
```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();\n\nranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=(tolink,(joined[1]/len(joined[0])))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter):\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .mapValues(lambda x: 0.15 + x * 0.85)
```



- The 2 RDDs are co-partitioned and co-located.
- Do you notice anything special??



Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ') .rdd.map(lambda x: (x[0], (x[1])))\n.distinct()\n.groupByKey()\n.cache();
```

```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=(tolink,(joined[1]/len(joined[0])))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter):\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .mapValues(lambda x: 0.15 + x * 0.85)
```

T-R	T-I	T-R I
R-B, T	R-I	R-B, T I

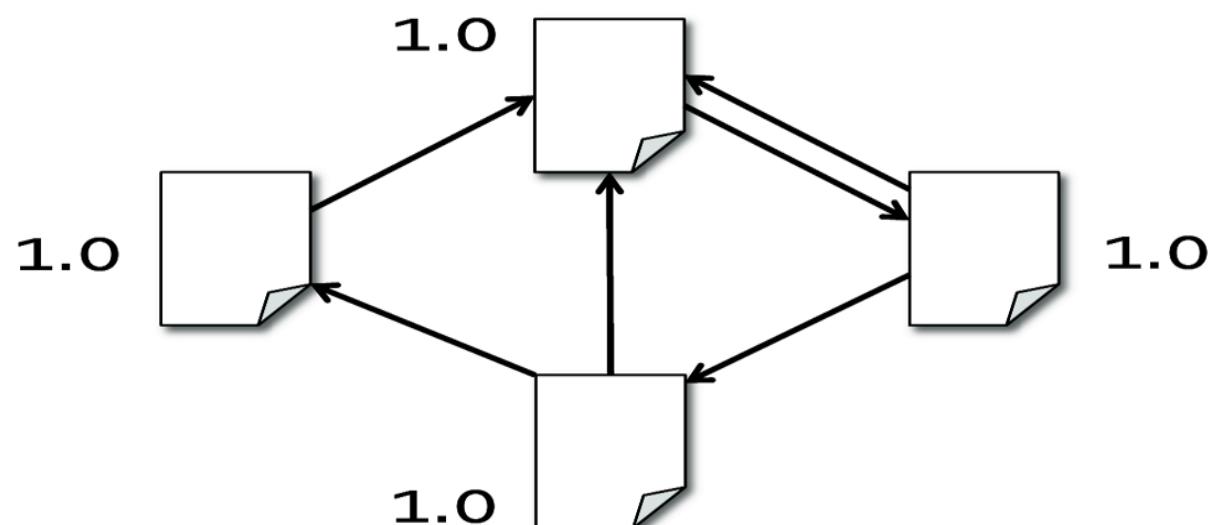
joined.values.

- No shuffle-sort for join

B-T, L	B-I	B-T, L I
L-T	L-I	L-T I

- The join function will be there **but no stage for join.**

Links Ranks joined



Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ').rdd.map(lambda x: (x[0], (x[1])))\n    .distinct()\n    .groupByKey()\n    .cache();
```

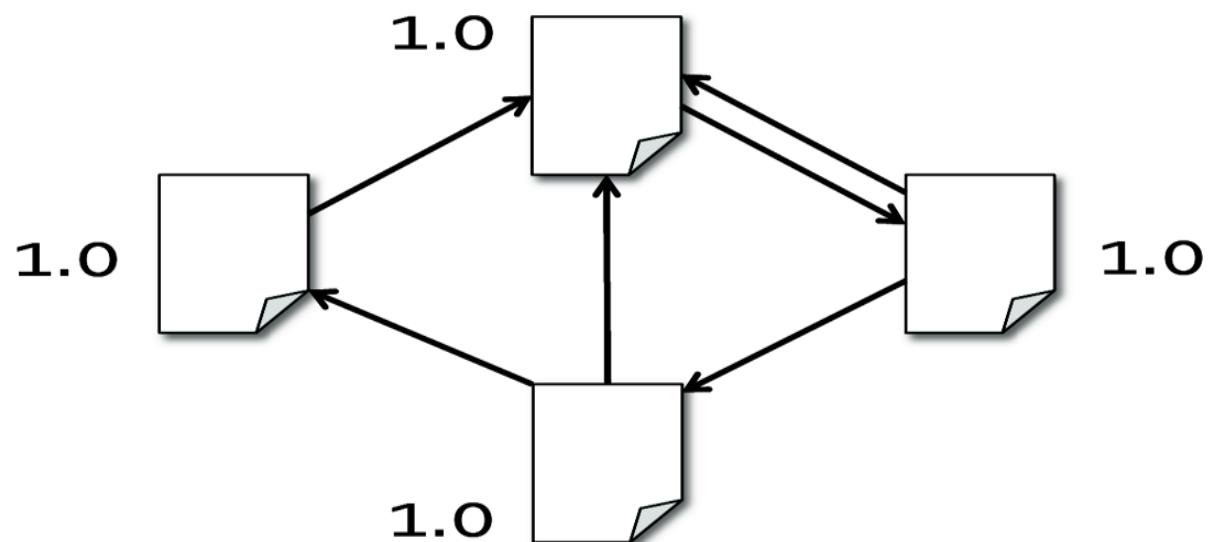
joined.values.

```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=(tolink,(joined[1]/len(joined[0])))\n        result.append(tup)\n    return result
```



```
for i in range(num_iter) :\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n        .flatMap(lambda joined: processJoined(joined))\n        .reduceByKey(add)\n        .mapValues(lambda x: 0.15 + x * 0.85)
```



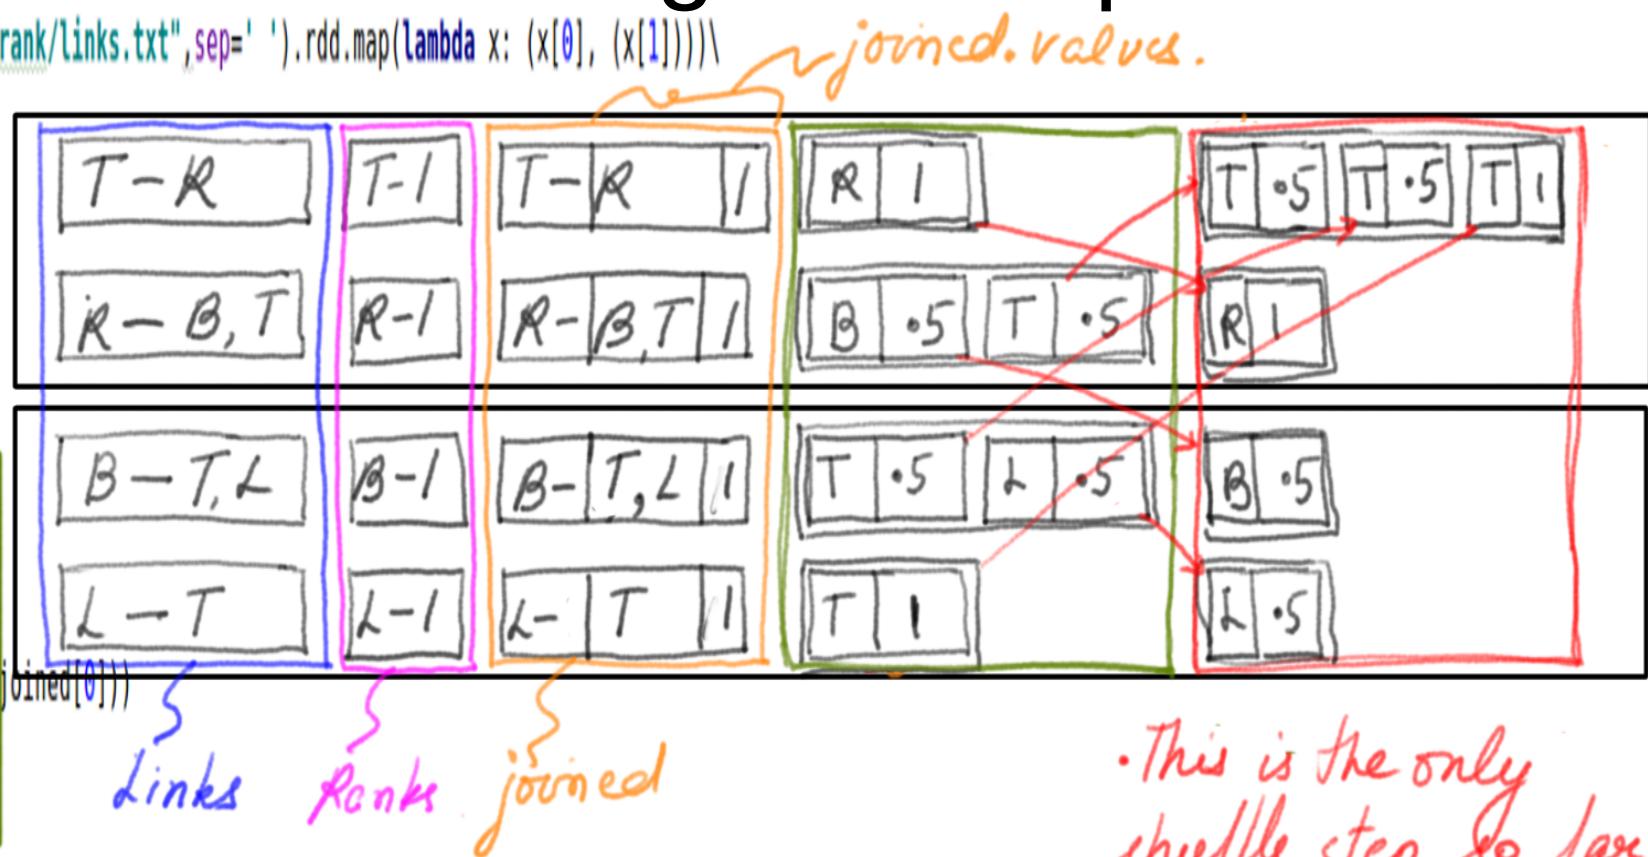
Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ')
.distinct()\n.groupByKey()\n.cache();
```

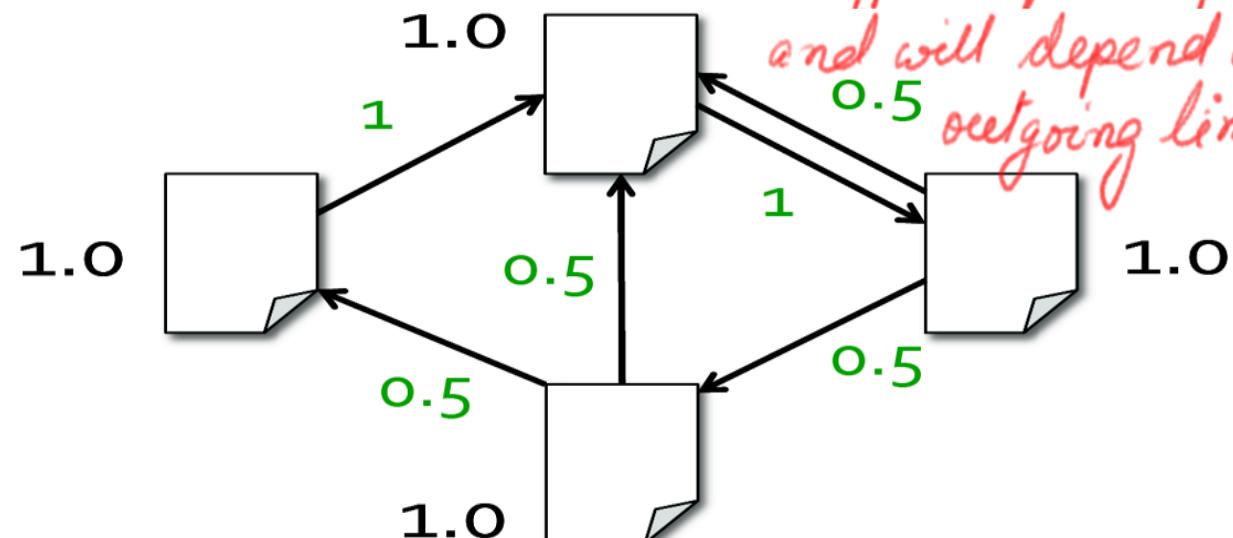
```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):\n    result=[]\n    for tolink in joined[0]:\n        tup=(tolink,(joined[1]/len(joined[0])))\n        result.append(tup)\n    return result
```

```
for i in range(num_iter):\n    #calculates URL contributions to the rank of other URLs.\n    ranks = links.join(ranks).values()\n    .flatMap(lambda joined: processJoined(joined))\n    .reduceByKey(add)\n    .mapValues(lambda x: 0.15 + x * 0.85)
```



This is the only shuffle step so far, and will depend on outgoing links.



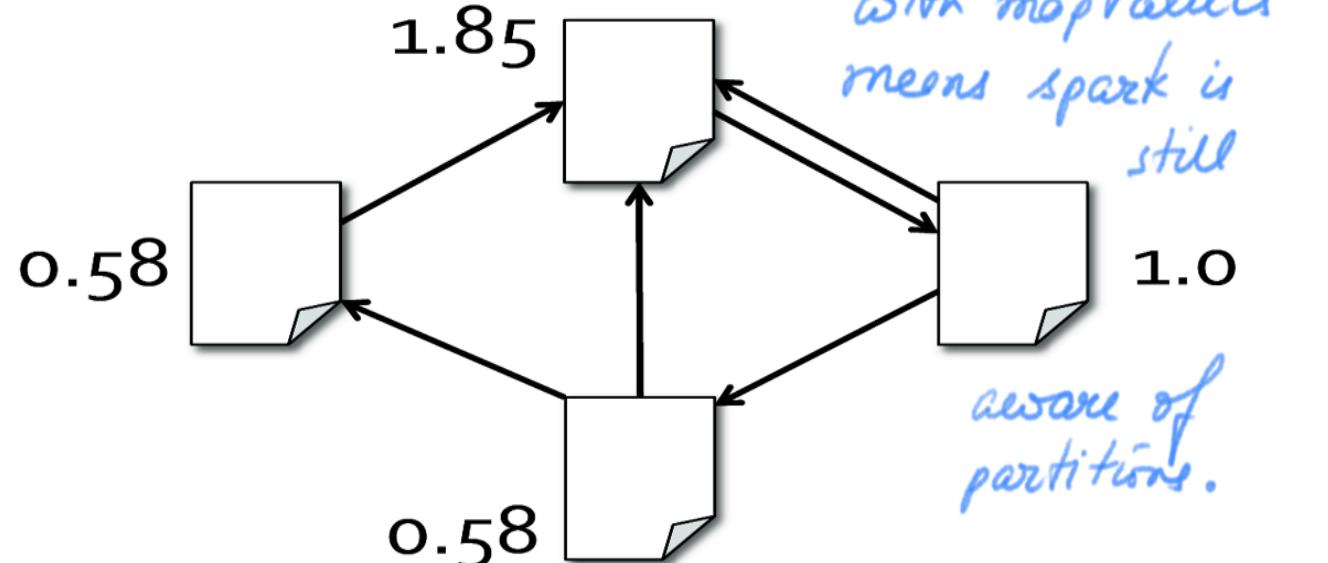
Spark Internals: Pagerank: Optimized

```
links = spark.read.csv("input/pagerank/links.txt", sep=' ')
    .distinct()
    .groupByKey()
    .cache();
```

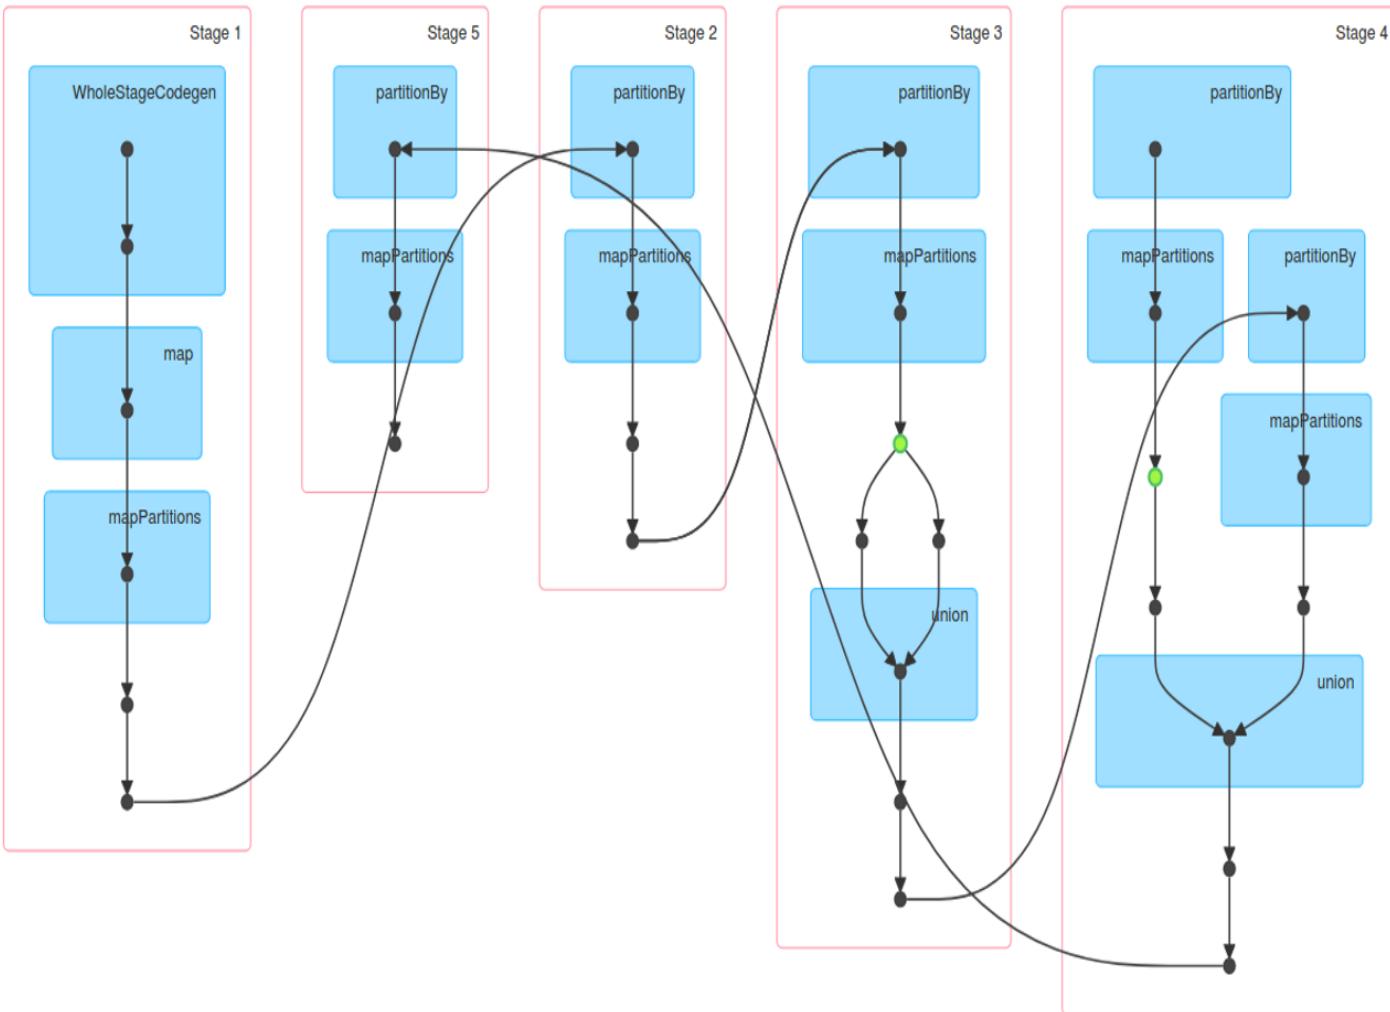
```
ranks=links.mapValues(lambda x:1)
```

```
def processJoined(joined):
    result=[]
    for tolink in joined[0]:
        tup=(tolink,(joined[1]/len(joined[0])))
        result.append(tup)
    return result
```

```
for i in range(num_iter):
    #calculates URL contributions to the rank of other URLs.
    ranks = links.join(ranks).values()
        .flatMap(lambda joined: processJoined(joined))
        .reduceByKey(add)
        .mapValues(lambda x: 0.15 + x * 0.85)
```



Spark Internals: Pagerank: Optimized

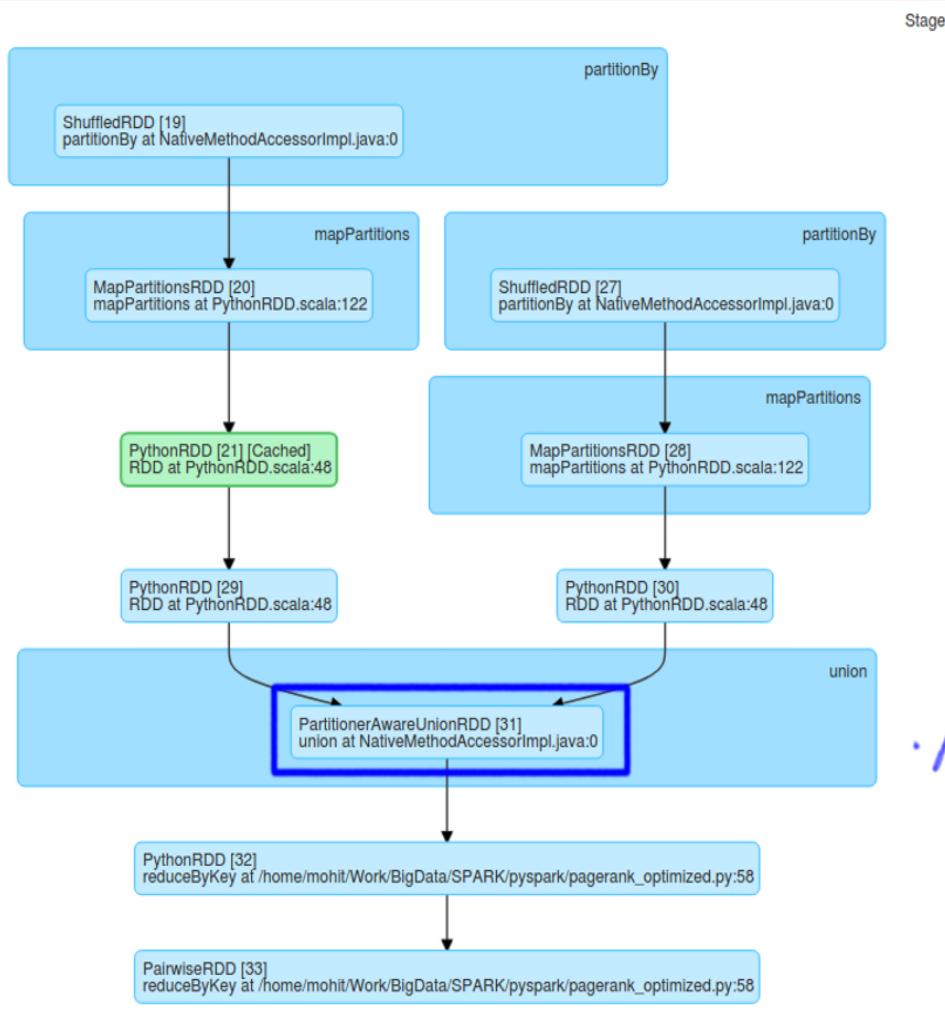


• No join stage, join function is there but not join stage as there is no shuffle.

Completed Stages (5)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	collect at /home/mohit/Work/BigData/SPARK/pyspark/pagerank_optimized.py:64	+details	2021/03/15 10:24:09	59 ms	1/1		169.0 B	
4	reduceByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank_optimized.py:58	+details	2021/03/15 10:24:09	0.1 s	1/1	310.0 B	163.0 B	169.0 B
3	reduceByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank_optimized.py:58	+details	2021/03/15 10:24:08	0.2 s	1/1	310.0 B	185.0 B	163.0 B
2	groupByKey at /home/mohit/Work/BigData/SPARK/pyspark/pagerank_optimized.py:33	+details	2021/03/15 10:24:08	69 ms	1/1		209.0 B	185.0 B
1	distinct at /home/mohit/Work/BigData/SPARK/pyspark/pagerank_optimized.py:33	+details	2021/03/15 10:24:08	0.4 s	1/1			209.0 B

Spark Internals: Pagerank: Optimized



Partition aware UnionRDD.

▶ Show Additional Metrics

▶ Event Timeline

Summary Metrics for 1 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0.1 s	0.1 s	0.1 s	0.1 s	0.1 s
GC Time	0 ms	0 ms	0 ms	0 ms	0 ms
Input Size / Records	310.0 B / 3	310.0 B / 3	310.0 B / 3	310.0 B / 3	310.0 B / 3
Shuffle Read Size / Records	163.0 B / 1	163.0 B / 1	163.0 B / 1	163.0 B / 1	163.0 B / 1
Shuffle Write Size / Records	169.0 B / 1	169.0 B / 1	169.0 B / 1	169.0 B / 1	169.0 B / 1